From	Johnny Lou <john...@campaignmonitor.com>
Subject	Re: ConsumerOffsets uneven partition distribution
Date	Wed, 21 Mar 2018 00:28:56 GMT

Hi Andras,

Thanks for that information, handcraft group.id to make sure it spread across broker is a way, I will give that a go.

I understand the benefit of consumer group, my concern at the moment is the potential to create a hot spot on one or the broker...

Thanks,

Johnny Luo

On 20/3/18, 10:03 pm, "Andras Beni" <andrasbeni@cloudera.com> wrote:

> Hi Johnny,
>
> As you already mentioned, it depends on the group.id which broker will be
> the group leader.
> You can change the group.id to modify which _consumer_offsets partition the
> group will belong to, thus change which broker will manage a group. You can
> check which partition a group.id is assigned using
>
> Utils.toPositive(Utils.murmur2(groupIdAsByteArray)) % partitionCount
>
> consumer group is a way to distribute work across equivalent consumers. I
> would assume it is a good idea but it depends on your architecture and use
> case.
>
> Best regards,
> Andras
>
> On Sat, Mar 17, 2018 at 12:55 PM, Johnny Luo <johnnyl@campaignmonitor.com>
> wrote:
>
> > Hello,
> >
> > We are running a 16 nodes kafka cluster on AWS, each node is a m4.xLarge
> > EC2 instance, with 2TB EBS(ST1) disk. Kafka version is 0.10.1.0, we have
> > about 100 topics at the moment. Some busy topics will have about 2 billion
> > events every day, some low volume topics will only have thousands per day.
> >
> > Most of our topics use an UUID as the partition key when we produce the
> > message, so the partitions are quite evenly distributed.
> >
> > We have quite a lot consumer consume from this cluster using consumer
> > group. Each consumer has a unique group id. Some consumer group commit
> > offsets every 500ms, some will commit offsets in sync as soon as it
> > finishes processing a batch of messages.
> >
> > Recently we observed a behaviour that some of the brokers are far busier
> > than the others. With some digging, we find out, it is actually quite a
> > lot traffic go to "__consumer_offsets", thus we created a tool to see the
> > high watermark of each partitions in "__consumer_offsets", which reveal
> > that the partitions are very uneven distributed.
> >
> > Based on this link "Consumer offset management in Kafka"
> >
> > It seems it is an intended behaviour, each consumer group only have one
> > leader, thus committed offsets all need to go to this leader, and also only
> > use “group.Id” to decide the partition.
> >
> > Given the fact that we have some consumers consume from those very busy

> topics, thus the commit offsets will cause a lot traffic to
> "__consumer_offsets" topic on the broker that handle the consumer group.
>
> My questions are :
> 1. Is there a way we can make sure that the consumer groups that consume
> from busy topics doesn't fall on to the same broker? Don’t' want to create
> a hotspot.
> 2. For consumers that consumer from busy topics (topics have billions
> messages per day), is it a good idea to use consumer group?
>
> Thanks in advance
>
> Johnny Luo
>

Mime	• Unnamed text/plain (inline, Quoted Printable, 3255 bytes) View raw message