
 [Answers \(/community/s/group/0F90L0000001NDBSA2/\)](/community/s/group/0F90L0000001NDBSA2/) — [Li wei \(/community/s/profile/0050L0000093E3HQAU\)](/community/s/profile/0050L0000093E3HQAU/) (Customer) asked a question.
[February 28, 2012 at 5:46 PM \(/community/s/question/0D50L00006BivL4SAL/mapr-file-system-io-performance-issue\)](/community/s/question/0D50L00006BivL4SAL/mapr-file-system-io-performance-issue)



MapR filesystem IO performance issue

I have just install M5 on 2 nodes ,each nodes hava this hardware configuration:

2*6 core

48G mem

2*10Gb NIC

14*15000rpm sas disk, 2 for RedhatLinux OS ,the rest 12 is for MapR,distsetup like :

`#disksetup -W 8 /home/disks`

`# more /home/disks`

`/dev/sdb`

`/dev/sdc`

`/dev/sdd`

`/dev/sde`

`/dev/sdf`

`/dev/sdg`

`/dev/sdh`

`/dev/sdi`

`/dev/sdj`

`/dev/sdk`

`/dev/sdl`

`/dev/sdm`

question is :

after created a volume on these two nodes, I do a dd test through NFS ,the result is bad:

`#mount vip:/mapr /hd` `#vip` is virtual IP for 2 NFS service which be started up in 2 two nodes;

`#cd /hd/MapR/mt` `#mt` is volume mount point

`#time dd if=/dev/zero of=2.test bs=32k count=4000000`

4000000+0 records in

4000000+0 records out

131072000000 bytes (131 GB) copied, 264.22 seconds, 496 MB/s

!!!just 496MB/s for two nodes, I had tested DD by EXT3 for each node (do 2 raid5 group for 12 disks in BIOS),and get a 1000MB/s for each node.

what's wrong with it ,how to improve MapR filesystem performance ?

[Mt \(/community/s/topic/0TO0L000004VOnVWAW/mt\)](/community/s/topic/0TO0L000004VOnVWAW/mt)

[Time \(/community/s/topic/0TO0L000004VOnWWAW/time\)](/community/s/topic/0TO0L000004VOnWWAW/time)

[Vip \(/community/s/topic/0TO0L000004VOnXWAW/vip\)](/community/s/topic/0TO0L000004VOnXWAW/vip)

+3 more


 Upvote

 Answer

 Share


8 answers

37 views

 [Ravindra Pindikura \(/community/s/profile/0050L0000093CQ7QAM\)](/community/s/profile/0050L0000093CQ7QAM/) (MAPR TECHNOLOGIES)
7 years ago

- Do you have nfs service running on only 1 node or both nodes ?

- If not, are you running dd on the same node that is running the nfs service ?



- When you run multiple dd, does it help if you fire them on both the nodes in parallel (different files) ?

Upvote Reply



Li wei (/community/s/profile/0050L0000093E3HQAU) (Customer)

7 years ago

▼

configuration for each node:

2 cpu 6 core,48G mem,2NIC with 10Gb ,14 sas disk with 15000rpm

Upvote Reply



Ravindra Pindikura (/community/s/profile/0050L0000093CQ7QAM) (MAPR TECHNOLOGIES)

7 years ago

▼

can you run a simple streaming test over hadoop ? There's a program we distribute called RWSpeedTest sitting inside maprfs-test-0.1.jar inside the /opt/mapr/hadoop/hadoop-0.20.0/lib directory ... it simply writes (or reads) a single file. To run it

1. set compression off in the dir where you will create the test file(s):

```
% hadoop fs -mkdir /a/b/c

% hadoop mfs -setcompression off /a/b/c
```

2. run the test -- it will print the MB/s obtained.

```
% hadoop jar lib/maprfs-test-01.jar com.mapr.fs.RWSpeedTest

usage: RWSpeedTest filename [-]megabytes uri

% hadoop jar lib/maprfs-test-01.jar com.mapr.fs.RWSpeedTest /a/b/c/testfile 1000 maprfs:///
```

You can run several copies of the RWSpeedTest in parallel to see what is aggregate b/w you get (run each one on a diff file)

Upvote Reply



Li wei (/community/s/profile/0050L0000093E3HQAU) (Customer)

7 years ago

▼

configuration for each node:

2*6 core,48G mem,2*10Gb NIC,14*15000rpm sas disk

Upvote Reply



Li wei (/community/s/profile/0050L0000093E3HQAU) (Customer)

7 years ago

▼

I use hdparm to test disk IO performance ,result can reach 196MB/s,

so it look like disk IO is ok ,but I found average is 5M~8M/s for earch disk by nmon monitor when I do dd test .

So bottlenech is not in disk , what's the bottleneck ,NFS ?

```
-----

hdparm -t /dev/sdc

/dev/sdc:

Timing buffered disk reads: 590 MB in 3.00 seconds = 196.40 MB/sec

[root@node1 ~]# hdparm -t /dev/sdm

/dev/sdm:

Timing buffered disk reads: 604 MB in 3.01 seconds = 200.95 MB/sec

[root@node1 ~]#

Upvote    Reply
```



Nabeel Moidu (/community/s/profile/005E0000000IT9aIAG) (MapR Technologies)


7 years ago

▼

What's the replication factor for your volumes?

Have you tried parallel writes to the volume ?

Upvote Reply

 Li wei (/community/s/profile/0050L0000093E3HQAU) (Customer) 7 years ago

replication factor =1;


I found result had no apparent difference for 1 or 3.

When I run 2 DD into same volume ,run time will double and 230M/s for each dd process:

```
#time dd if=/dev/zero of=2.test bs=32k count=4000000 &

#time dd if=/dev/zero of=3.test bs=32k count=4000000 &

Upvote   Reply
```

 Li wei (/community/s/profile/0050L0000093E3HQAU) (Customer) 7 years ago

I have just install M5 on 2 nodes ,each nodes hava this hardware configuration:

2*6 core

48G mem

2*10Gb NIC

14*15000rpm sas disk, 2 for RedhatLinux OS ,the rest 12 is for MapR,distsetup like :

```
#disksetup -W 8 /home/disks

# more /home/disks

/dev/sdb

/dev/sdc

/dev/sdd

/dev/sde

/dev/sdf

/dev/sdg

/dev/sdh

/dev/sdi

/dev/sdj

/dev/sdk

/dev/sdl

/dev/sdm
```

question is :

after created a volume on these two nodes, I do a dd test through NFS ,the result is bad:

```
#mount vip:/mapr /hd #vip is virtual IP for 2 NFS service which be started up in 2 two nodes;
```

```
#cd /hd/MapR/mt #mt is volume mount point
```

```
#time dd if=/dev/zero of=2.test bs=32k count=4000000

4000000+0 records in

4000000+0 records out

13107200000 bytes (131 GB) copied, 264.22 seconds, 496 MB/s
```

!!!just 496MB/s for two nodes, I had tested DD by EXT3 for each node (do 2 raid5 group for 12 disks in BIOS),and get a 1000MB/s for each node.

what's wrong with it ,how to improve MapR filesystem performance ?

Upvote Reply



Write an answer...