

[MAPR USE-CASES](#) > [MAPR SECURITY](#) > [AUDIT](#)
(/support/s/login)

Why am I getting duplicate FID entries for single file operations in audit logs? Is coalesce interval functionality not working?

For every single "write" operation and for each chunk created, there is an audit record written to the audit logs. If you create a big file (e.g. 10GB), there will be many audit records (e.g. ~40) for the same FID. As these logs have the same UID, IP, FID, operation type, the audit-logging appears to be duplicated (although coalescing is enabled and the operations are executed within the same millisecond). This behavior is current and is as designed.

🕒 Jan 26, 2018 · Known Issue and EBF

Title

Why am I getting duplicate FID entries for single file operations in audit logs? Is coalesce interval functionality not working?

Article Type

Known_Issue_EBFs__kav

Article Number

000003158

MapR or Open Source Bug ID ⓘ

30408

Primary Open Source Bug Link ⓘ

Dependent Open Source Bug Link(s)

None

Bug Description ⓘ

For a single operation on any individual file in a volume, there appear to be multiple or duplicate entries in the fileserver audit logs.

To reproduce the behavior:

1. Create a volume and set audit on the mounted directory as well as on volume.

```
# maprcli volume create -name testvolume -path /testvolume -auditenabled true

# hadoop mfs -setaudit on /testvolume

# maprcli volume info -name testvolume -json | grep -i 'audited\|coalesce'

"audited":1,

"coalesceInterval":60,
```

2. Create a 1GB file using NFS/FUSE or copy a large file using the hadoop shell.

```
# dd if=/dev/zero of=/mapr/my.cluster.com/testvolume/filename (http://my.cluster.com/testvolume/filename) bs=1024 count=1048576

1048576+0 records in

1048576+0 records out

1073741824 bytes (1.1 GB) copied, 5.65231 s, 190 MB/s
```

```
# hadoop mfs -ls /testvolume

Found 1 items

-rw-r--r--  Z U A   3 root root 1073741824 2018-01-10 19:23 268435456 /testvolume/filename

    p 2696.32.131262  node3:5660 node2:5660 mapr-hadoop4:5660

    0 2697.32.131250  node2:5660 mapr-hadoop4:5660 node3:5660

    1 2699.32.131272  node2:5660 mapr-hadoop4:5660 node3:5660

    2 2700.32.131216  node2:5660 node3:5660 mapr-hadoop4:5660

    3 2701.32.131358  node2:5660 mapr-hadoop4:5660 node3:5660
```

3. Check the fileserver audit logs for the nodes where the file chunks are written. 'node2' in this case per the above output.

```
# hadoop fs -cat /var/mapr/local/node2/audit/5660/FSAudit.log-2018-01-10-001.json | grep WRITE
{"timestamp":{"$date":"2018-01-10T23:30:10.954Z"},"operation":"WRITE","uid":0,"ipAddress":"10.10.73.26","nfsServer":"10.10.73.26"}
{"timestamp":{"$date":"2018-01-10T23:30:12.537Z"},"operation":"WRITE","uid":0,"ipAddress":"10.10.73.26","nfsServer":"10.10.73.26"}
{"timestamp":{"$date":"2018-01-10T23:30:14.061Z"},"operation":"WRITE","uid":0,"ipAddress":"10.10.73.26","nfsServer":"10.10.73.26"}
{"timestamp":{"$date":"2018-01-10T23:30:15.747Z"},"operation":"WRITE","uid":0,"ipAddress":"10.10.73.26","nfsServer":"10.10.73.26"}
```

Observation:

- For every single "write" operation and for each chunk created, there is an audit record written to the audit logs.
- If you create a big file (e.g. 10GB), there will be many audit records (e.g. ~40) for the same FID.
- As these logs have the same UID, IP, FID, operation type, the audit-logging appears to be duplicated (although coalescing is enabled and the operations are executed within the same millisecond).

Bug Found in Version? ⓘ

MapR 5.x,6.x

Work Around ⓘ

No workaround at this time.

Environment

Root Cause

The current behavior of the audit logs is that MapR audits once for each FID/chunk and records only the primary FID rather than actual FID. The rationale for recording primary FID instead of multiple FIDs for the same file is to eliminate the possible confusion of a user treating each FID as a separate/unique file. In effect, the 'primary FID' concept generates less confusion.

Bug Fixed in Version?

No fix at this time.

Technology Group

MapR Core

Article Total View Count

52

Last Published Date

1/26/2018 1:23 AM
