



Apache Drill – Evolution, Use cases & Roadmap

Neeraja Rentachintala, Senior Director of Product Management

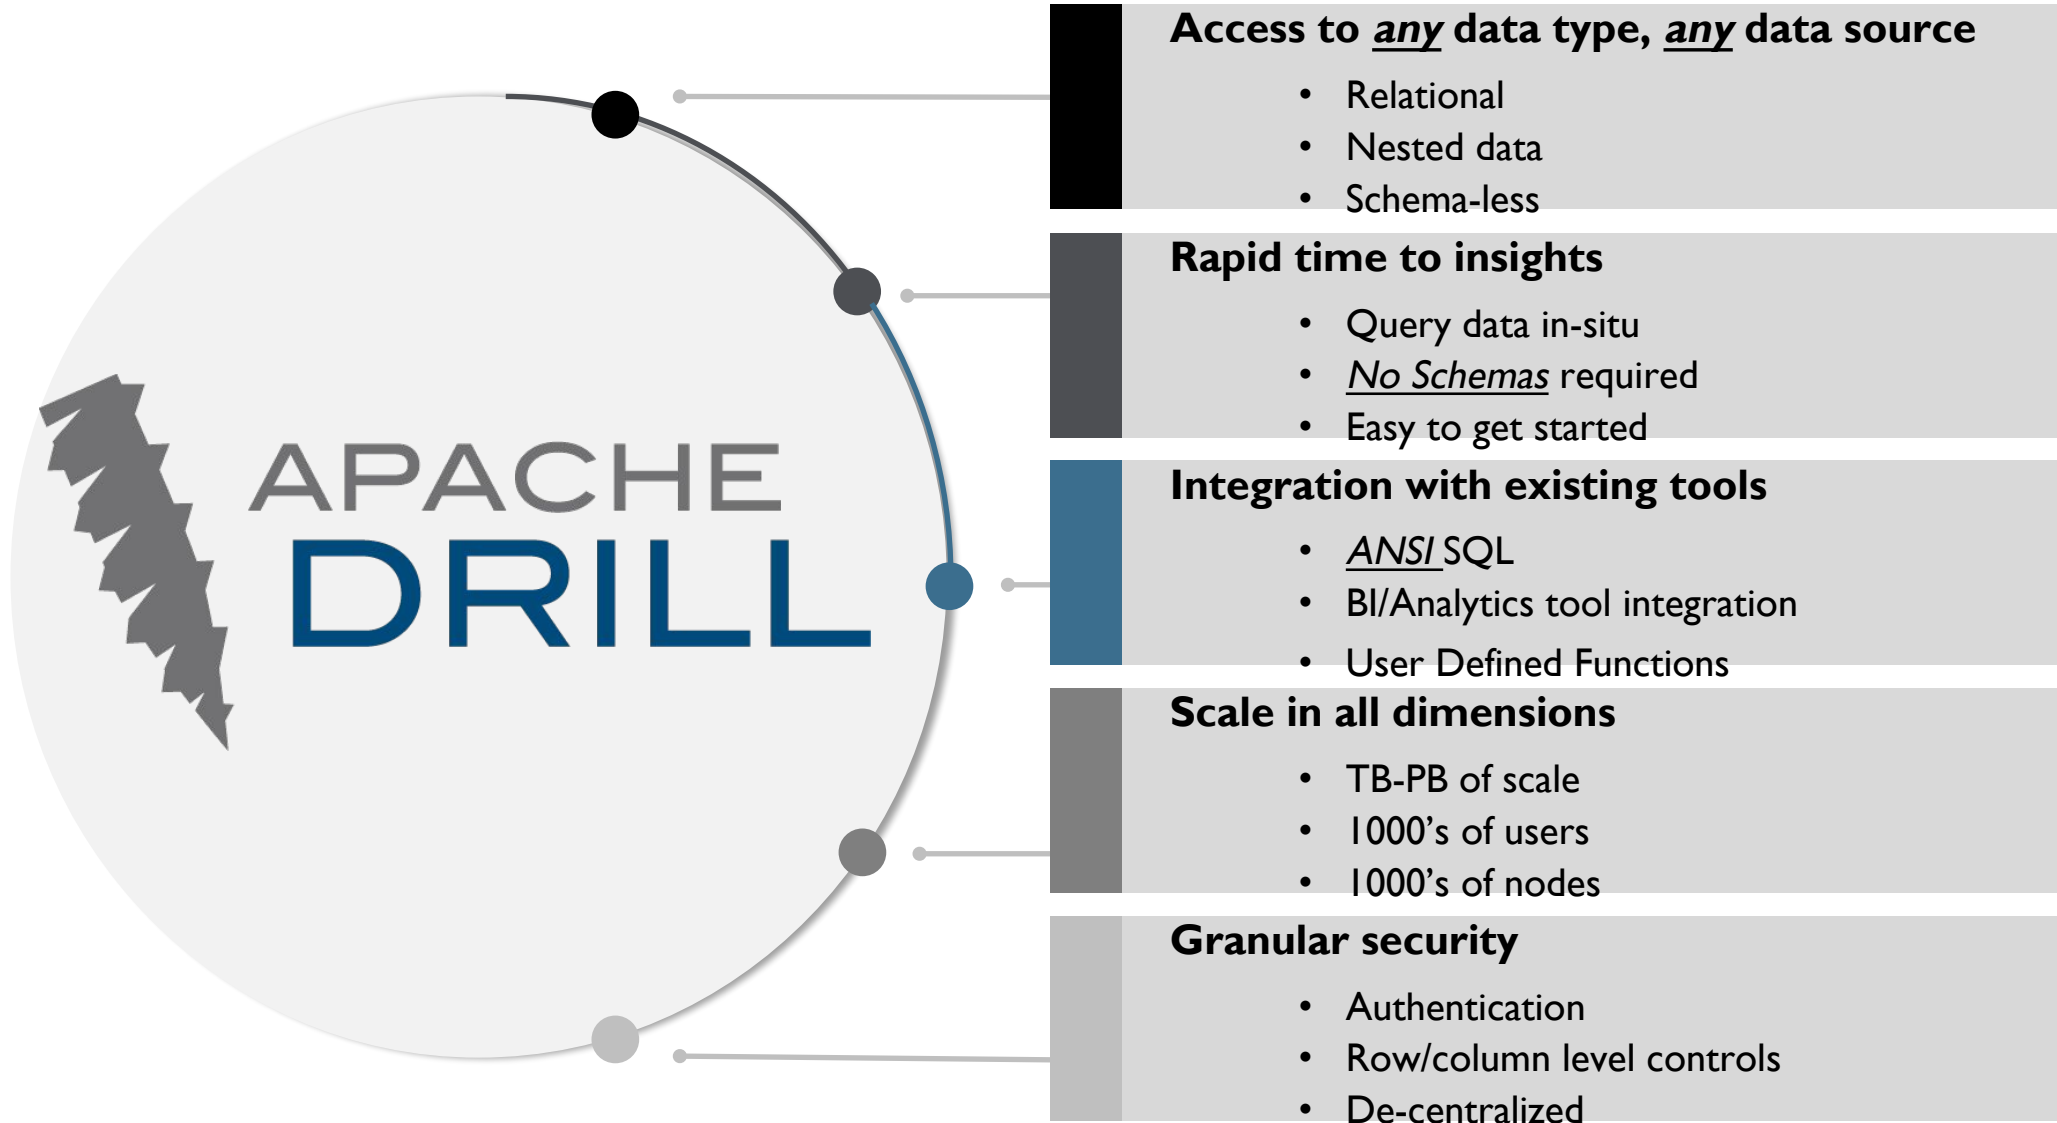
12/6/2016



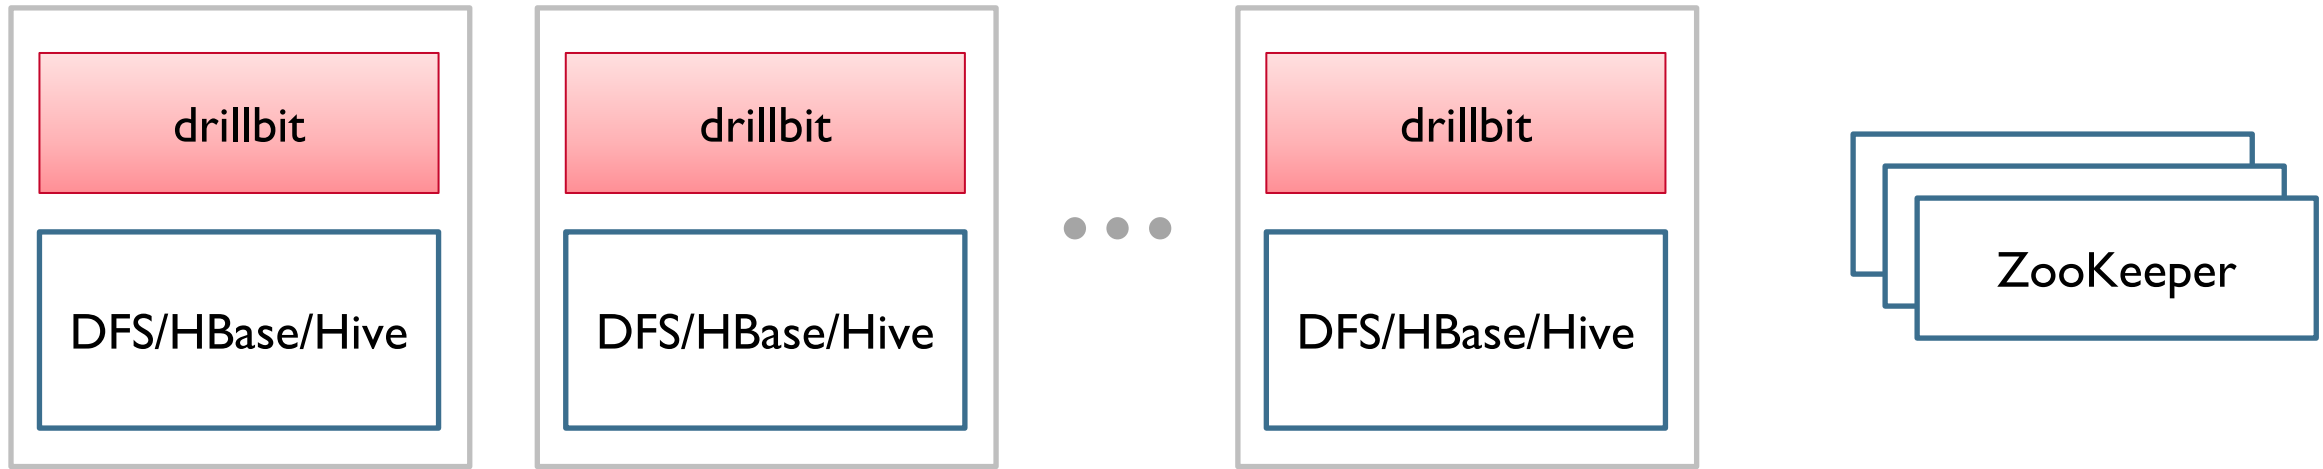
Topics

- Apache Drill overview
- Use cases
- Product evolution & roadmap

Industry's First Schema-free SQL engine for Big Data



Drill is a distributed query engine



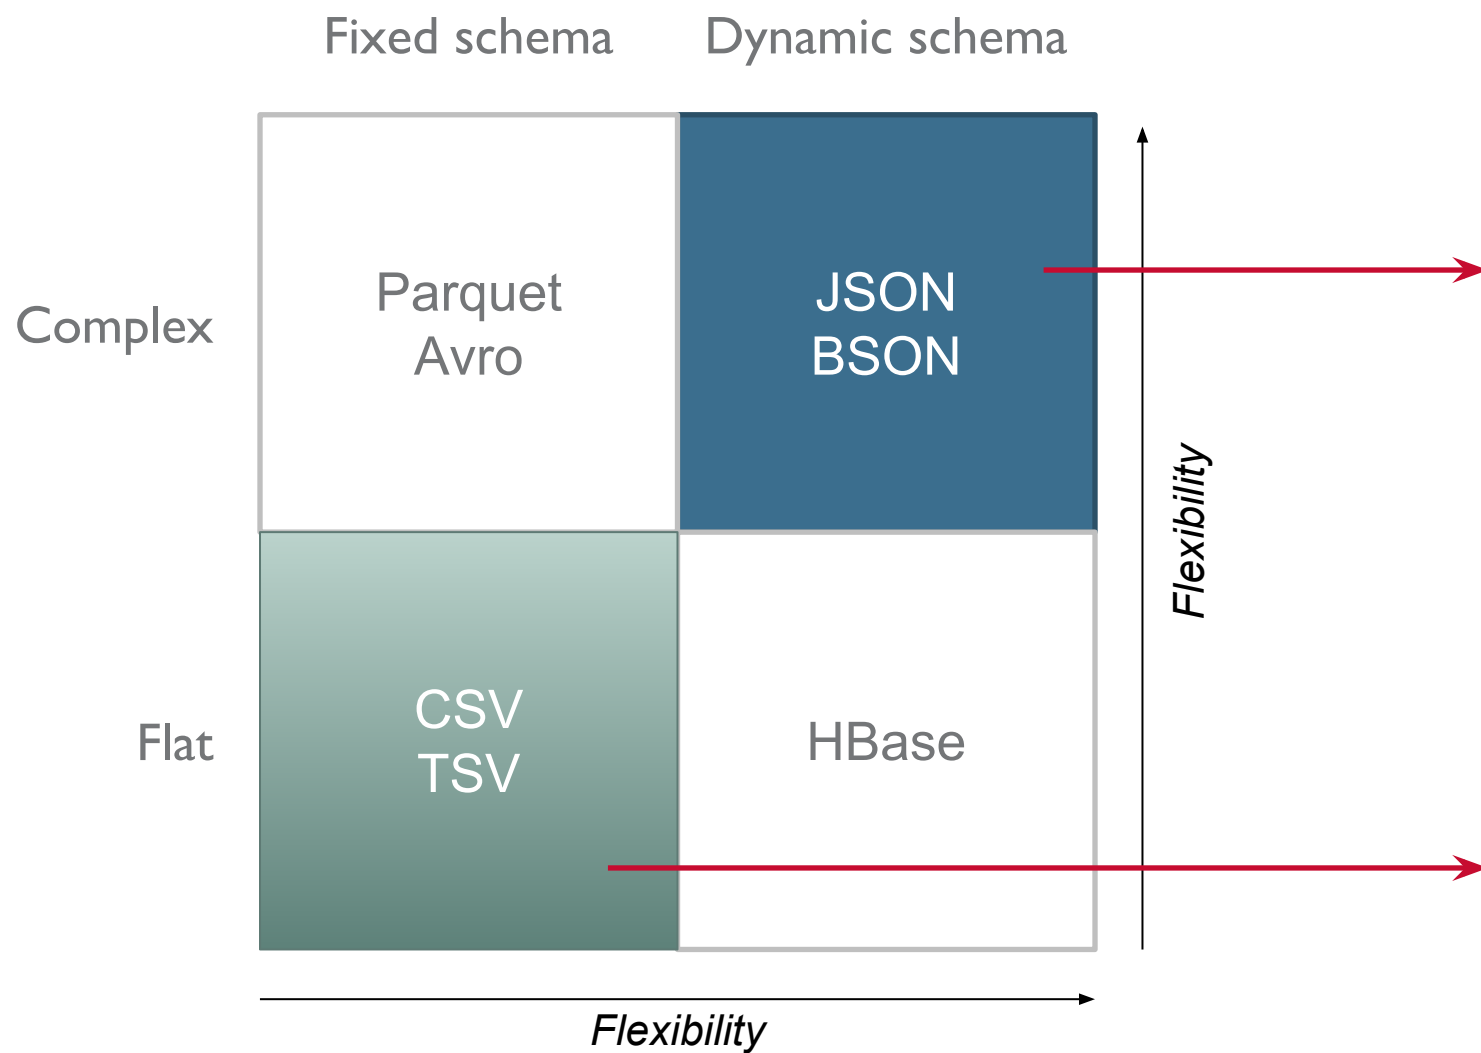
➤ Execution in Drill

- Scale-out MPP
- Hierarchical “JSON like” data model
- Columnar and Vectorized processing
- Optimistic execution (no MR, Spark)
- Runtime compilation
- Late binding
- Extensible

➤ Optimization in Drill

- Apache Calcite+ Parallel optimizations
- Data locality awareness
- Projection pruning
- Filter pushdown
- Partition pruning
- CBO & pluggable optimization rules
- Metadata caching

Drill's Data Model is Flexible



Apache Drill table

```
{
  name: {
    first: Michael,
    last: Smith
  },
  hobbies: [ski, soccer],
  district: Los Altos
}
{
  name: {
    first: Jennifer,
    last: Gates
  },
  hobbies: [sing],
  preschool: CCLC
}
```

RDBMS/SQL-on-Hadoop table

Name	Gender	Age
Michael	M	6
Jennifer	F	3



Enabling “As-It-Happens” Business with Instant Analytics

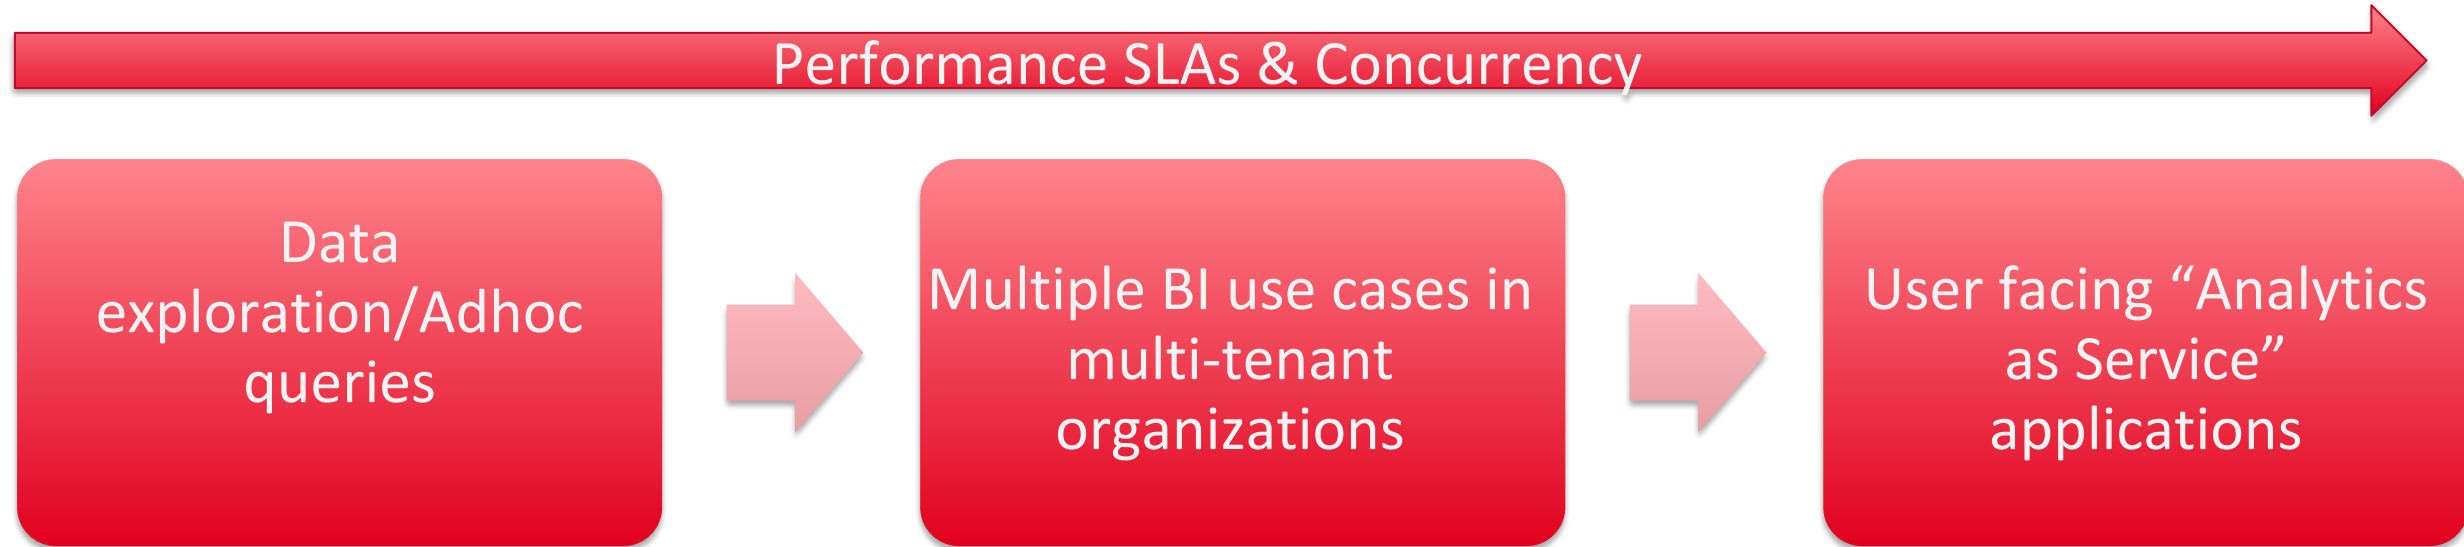
Traditional approach



Exploratory approach with Drill



Drill Enables Traditional and New types of Analytics @ Scale



Example - Data exploration/Adhoc queries

Lets take a IOT race cars scenario

Producers



Sensors data



Sensor Data VI

```
{
  "_id" : "car1-1.462271134E9/102.982",
  "car" : "car1",
  "race_id" : "20160503102534",
  "racetime" : 102.982,
  "records" : [ {
    "race_id" : "20160503102534",
    "racetime" : 102.982,
    "sensors" : {
      "Distance" : 1107.563599,
      "RPM" : 1241.948975,
      "Speed" : 29.808878
    },
    "timestamp" : 1.462271134E9
  }, {
    "race_id" : "20160503102534",
    "racetime" : 103.172,
    "sensors" : {
      "Distance" : 1113.882324,
      "RPM" : 1244.506958,
      "Speed" : 29.794189
    },
    "timestamp" : 1.462271134E9
  }
}
```


Namespace of a Drill query

Workspace

- Sub-directory
- HBase namespace
- Hive database

Table

- Pathnames
- HBase table
- Hive table

```
SELECT * FROM dfs.demo.`sample_car_data_v1.json`
```

Storage plugin instance

- File system (Text, Parquet, JSON)
- HBase/MapR-DB
- Hive Metastore/HCatalog
- Easy API to go beyond Hadoop (MongoDB, JDBC, ES etc)

Drill has ability to discover schemas on the fly

```
$ tar -xvzf apache-drill-1.9.0.tar.gz
```

```
$ bin/sqlline -u jdbc:drill:zk=local
```

```
> SELECT * FROM dfs.demo.`sample_car_data_v1.json` LIMIT 1;
```

```
+-----+-----+-----+-----+-----+-----+
|_id | car | race_id | racetime | records | timestamp |
+-----+-----+-----+-----+-----+-----+
| car1-1.462271134E9/102.982 | car1 | 20160503102534 | 102.982 |
[{"race_id":"20160503102534","racetime":102.982,"sensors":{"Distance":1107.563599,"RPM":1241.948975,"Speed":29.808878},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":103.172,"sensors":{"Distance":1113.882324,"RPM":1244.506958,"Speed":29.794189},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":103.322,"sensors":{"Distance":1118.141846,"RPM":1249.089233,"Speed":29.776337},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":103.572,"sensors":{"Distance":1126.747314,"RPM":1243.363037,"Speed":29.756483},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":103.786,"sensors":{"Distance":1133.979004,"RPM":1240.901978,"Speed":29.732235},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":103.958,"sensors":{"Distance":1139.061035,"RPM":1243.04248,"Speed":29.699705},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":104.138,"sensors":{"Distance":1145.604126,"RPM":1241.426025,"Speed":29.664589},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":104.3,"sensors":{"Distance":1150.693237,"RPM":1237.156738,"Speed":29.629286},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":104.546,"sensors":{"Distance":1158.682251,"RPM":1235.059326,"Speed":29.567074},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":104.722,"sensors":{"Distance":1164.483276,"RPM":1233.207031,"Speed":29.520657},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":104.878,"sensors":{"Distance":1169.551392,"RPM":1230.01062,"Speed":29.481007},"timestamp":1.462271134E9},{"race_id":"20160503102534","racetime":112.658,"sensors":{"Distance":1408.384399,"RPM":1251.863525,"Speed":29.924427},"timestamp":1.462271134E9}] | 1.462271134E9 |
+-----+-----+-----+-----+-----+-----+
```

Query and Manipulate complex data (Contd)

```
> SELECT
    car,
    AVG(t.records.sensors.Speed) AS `m/s`,
    (AVG(t.records.sensors.Speed)*18)/5 AS `km/h`
FROM
    (SELECT *,FLATTEN(records) AS records
     FROM dfs.demo.`sample_car_data_v1.json`) t
GROUP BY car;
```

car	m/s	km/h
car1	48.712233453333305	175.36404043199988
car2	32.86935792448977	118.32968852816316
car3	46.06944574676411	165.8500046883508
car4	45.020420319341206	162.07351314962835
car5	44.31281895333338	159.52614823200014
car6	44.57982400883707	160.48736643181346
car7	45.07025832114878	162.2529299561356
car8	44.91297666539465	161.68671599542074
car9	43.887101968333305	157.9935670859999

Sensor Data V2

- 5 main data points:
 - Speed (m/s)
 - RPM
 - Distance (m)
 - **Fuel**
 - **Gear**



New fields added
new sensors

```
{
  "_id" : "car3-1.462279145E9/2.342",
  "car" : "car3",
  "race_id" : "20160503123905",
  "racetime" : 2.342,
  "records" : [ {
    "race_id" : "20160503123905",
    "racetime" : 2.342,
    "sensors" : {
      "Distance" : 2011.013672,
      "Fuel" : 4.843411,
      "Gear" : 2.0,
      "RPM" : 611.679199,
      "Speed" : 22.373652
    }
  }, {
    "race_id" : "20160503123905",
    "racetime" : 3.176,
    "sensors" : {
      "Distance" : 2031.738525,
      "Fuel" : 4.822249,
      "Gear" : 2.0,
      "RPM" : 722.421448,
      "Speed" : 26.686924
    }
  }
],
  "timestamp" : 1.462279145E9
}
```

Query on new fields added in V2 instantly

```
> SELECT
    car,
    t.records.sensors.gear AS gear,
    AVG(t.records.sensors.Speed) AS `m/s`,
    (AVG(t.records.sensors.Speed)*18)/5 AS `km/h`
FROM
    (SELECT *, FLATTEN(records) AS records
     FROM dfs.demo.`sample_car_data_v2.json`) t
GROUP BY car , t.records.sensors.gear
ORDER BY car , t.records.sensors.gear LIMIT 10;
```

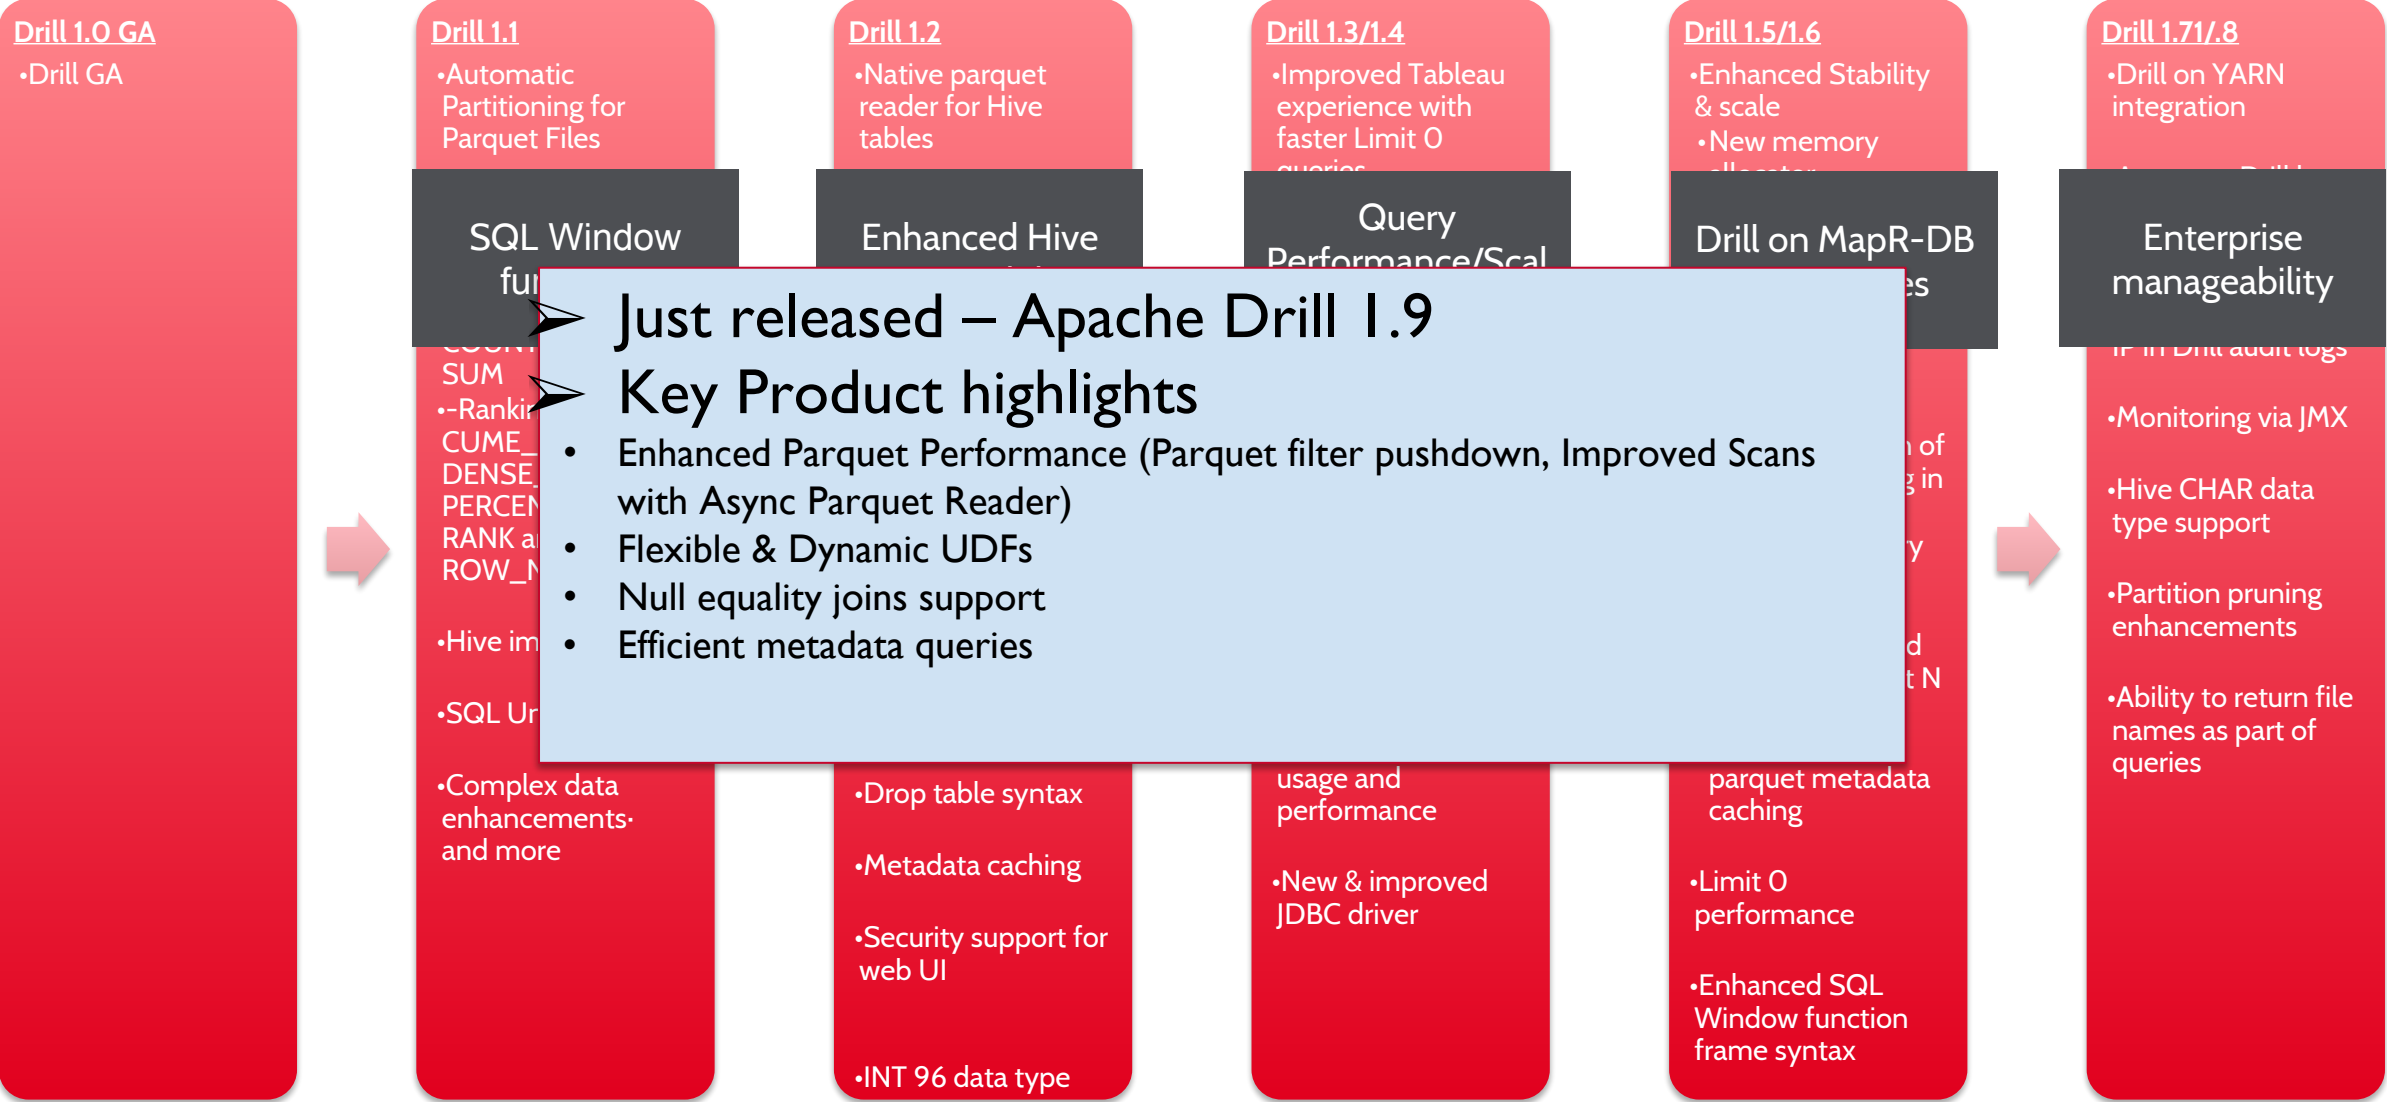
car	gear	m/s	km/h
car3	1.0	11.897254207792207	42.830115148051945
car3	2.0	27.522497517799366	99.08099106407772
car3	3.0	37.204457530487815	133.93604710975615
car3	4.0	49.01672581017499	176.46021291662996
car3	5.0	61.91693234567113	222.90095644441607
car3	6.0	70.04746468410457	252.17087286277646
car4	1.0	10.57082569811321	38.05497251320755
car4	2.0	24.364864507246377	87.71351222608696
car4	3.0	36.6934715200831	132.09649747229918
car4	4.0	50.961303641552576	183.46069310958927

Drill provides ANSI SQL capabilities

```
> WITH X AS (  
    SELECT car,  
    t.race_id AS race,  
    t.records.sensors.speed AS speed ,  
    RANK() OVER (PARTITION BY t.car, t.race_id  
        ORDER BY t.records.sensors.speed DESC) AS carracerank  
    FROM (SELECT *, FLATTEN(records) AS records FROM  
        dfs.demo.`sample_car_data_v1.json`) t)  
SELECT X.Car, X.race, X.speed FROM X  
WHERE X.carracerank =1 LIMIT 15;
```

Car	race	speed
car1	20160503102534	83.075401
car2	20160503102534	48.842308
car2	20160503102837	45.507034
car3	20160503103501	67.85479
car3	20160503103937	67.804237
car3	20160503104300	67.866081
car4	20160503102534	68.376411
car4	20160503102837	63.95813
car4	20160503103501	68.387604
car4	20160503103937	68.398613
car4	20160503104300	68.417183
car5	20160503102534	68.214432
car6	20160503102534	68.877663
car6	20160503102837	64.2967
car6	20160503103501	70.95742

Drill product evolved significantly since GA



Simplified deployment with YARN

- Drill as a long running application in YARN
- Key features
 - Client tool to launch Drill as YARN application
 - New Drill application master (AM)
 - CPU & memory controls
 - Add/remove nodes to cluster
 - Multiple Drill clusters

Apache Drill Configuration Drillbits Manage History						
YARN Application Master – Drill-on-YARN						
Drill Cluster Status						
YARN Application ID:	application_1465323848842_0024					
YARN Resource Manager:	doy1.mapr.lab					
YARN Node Manager for AM:	doy8.mapr.lab App info					
ZooKeeper:	10.10.102.25:5181/drill/drillbits1					
State:	LIVE					
Target Drillbit Count:	8					
Live Drillbit Count:	6					
Total Drill Virtual Cores:	12					
Total Drill Memory (MB):	36,864					
Yarn Node Count:	7					
Available Yarn Virtual Cores:	14					
Available Yarn Memory (MB):	8,192					
Pool	Type	Target Drillbit Count	Total Drillbits	Live Drillbits	Memory per Drillbit (MB)	VCores per Drillbit
pool-1	basic	8	8	6	6,144	2

Drill configuration w/YARN

Drill on YARN (contd)

Drill cluster management under YARN

Apache Drill Configuration Drillbits Man

Manage Drill Cluster

Current Status: 2 Drillbits are running.
Free YARN nodes: Approximately 6

Resize to drillbits.

Stop the Drill cluster.

Drill cluster status under YARN

YARN Application Master – Drill-on-YARN					
Drillbit Status					
ID	Pool	Host	State	ZK State	Container
1	pool-1	doy1.mapr.lab	RUNNING [x]	START_ACK	container_e08_1465323848842_0024_01_000002
2	pool-1	doy2.mapr.lab	RUNNING [x]	START_ACK	container_e08_1465323848842_0024_01_000003
3	pool-1	doy5.mapr.lab	RUNNING [x]	START_ACK	container_e08_1465323848842_0024_01_000004
4	pool-1	doy7.mapr.lab	RUNNING [x]	START_ACK	container_e08_1465323848842_0024_01_000005
5	pool-1	doy8.mapr.lab	RUNNING [x]	START_ACK	container_e08_1465323848842_0024_01_000006
6	pool-1	doy4.mapr.lab	RUNNING [x]	START_ACK	container_e08_1465323848842_0024_01_000007
7	pool-1		REQUESTING [x]	NEW	
8	pool-1		REQUESTING [x]	NEW	

New JMX based monitoring

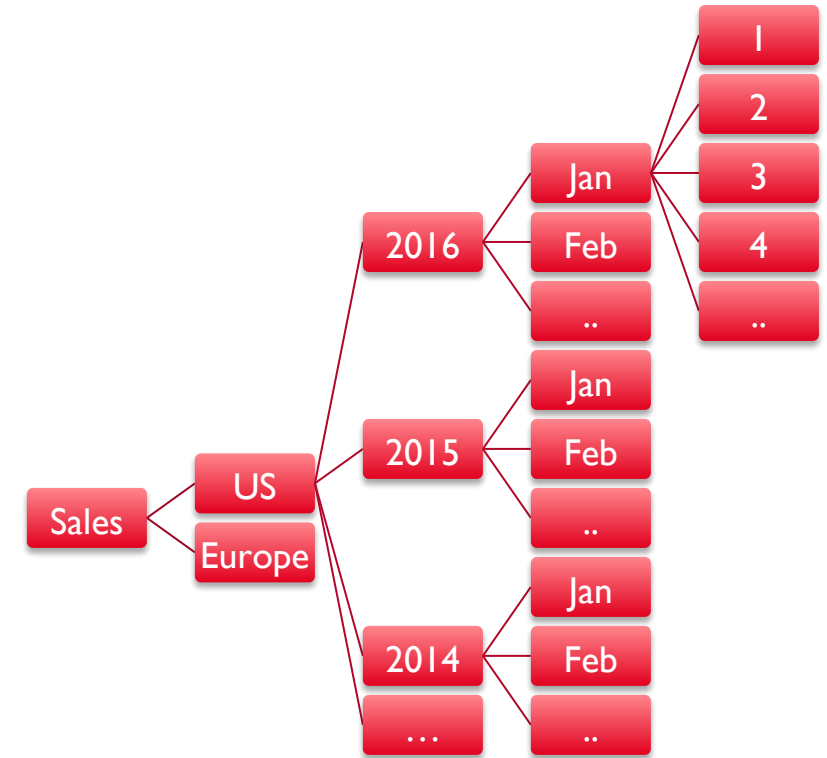
- New JMX based metrics Drill Web Console or MapR Spyglass (Beta) or a remote JMX monitoring tool, such as Jconsole
- Various system and query metrics
 - drill.queries.running
 - drill.queries.completed
 - heap.used
 - direct.used
 - waiting.count ...



Drill metrics dashboard with Spyglass

Improved query performance with Partition pruning

- Partition pruning allows a query engine to determine and retrieve the smallest needed dataset to answer a given query
- Data can be partitioned
 - At the time of ingestion into the cluster
 - As part of ETL via Hive or Spark or other batch processing tools
 - Drill support CTAS with PARTITION BY clause
- Drill does partition pruning for queries on partitioned Hive tables as well as file system queries



Select * from Sales
Where dir0='US' and dir1 ='2015'

Metadata Caches to Speed up Query Planning

- Helps reduce query planning time significantly when working with large # of Parquet files (thousands to millions)
- Highly optimized cache with the key metadata from parquet files
 - Column names, data types, nullability, row group size...
- Recursive cache creation at root level or selectively for specific directories or files
 - Ex: `REFRESH TABLE METADATA dfs.tmp.BusinessParquet;`
- Metadata caching is better suited for large amounts of data with moderate rate of change

Improved Query performance with Parquet Filter pushdown

- Applies during planning time
- Evaluates filter condition before the scan
- Planner evaluates filter conditions and checks if a Parquet row group can be eliminated
- Requires Parquet files to have min/max statistics

Example

```
SELECT * from table_t1  
WHERE      date_column  between      date  
           '2016-01-01' and date '2016-01-31'
```


Row group 1 : date_column : min =
2015-01-01 max = 2015-12-31

Row group 2 : date_column : min =
2016-01-01 max = 2016-12-31

Only row group 2 will be scanned

Parquet Filter pushdown

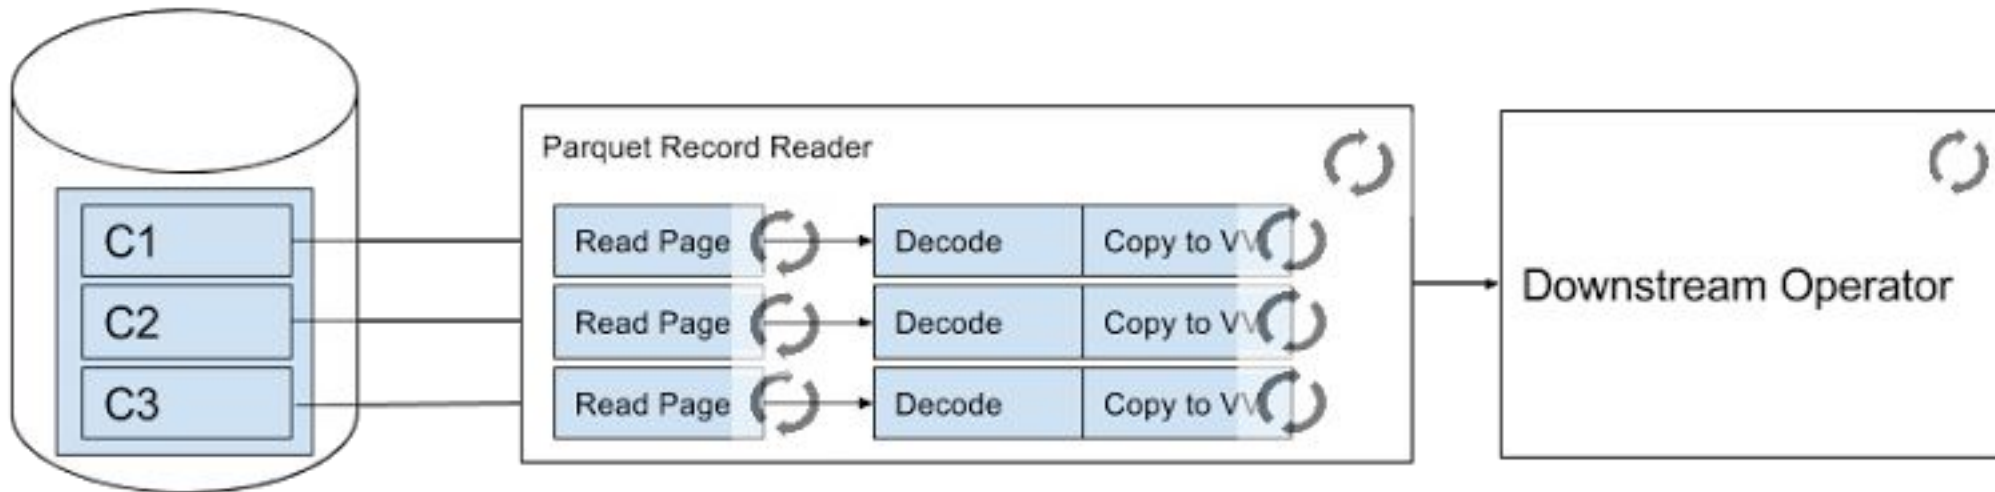
Parquet Filter Pushdown (Ver 1)
TPC-H SF1000 - Single User



TPCH Queries	Selectivity	Without FltrPD (MB)	With FltrPD (MB)	I/O Reduction
TPCH 06	15%	5,779	1,707	70%
TPCH 07	30%	12,395	5,188	58%
TPCH 14	1%	7,915	5,254	34%
TPCH 20	15%	9,174	8,333	9%

Enhanced query performance w/Asynchronous Parquet reader

- High performance queries for scan intensive analytics (~33% I/O reduction)
- Parquet reader improvements include
 - Buffered reads
 - Parallel reads from file system
 - Parallel decompression and decoding
 - Reading and decoding is pipelined



How all these fit together?

Hive/Spark on
File System
data
(Batch/ETL
workloads)



Drill + FS w/raw file
formats (ex: Text,
JSON..)
(Data exploration/
Adhoc queries)

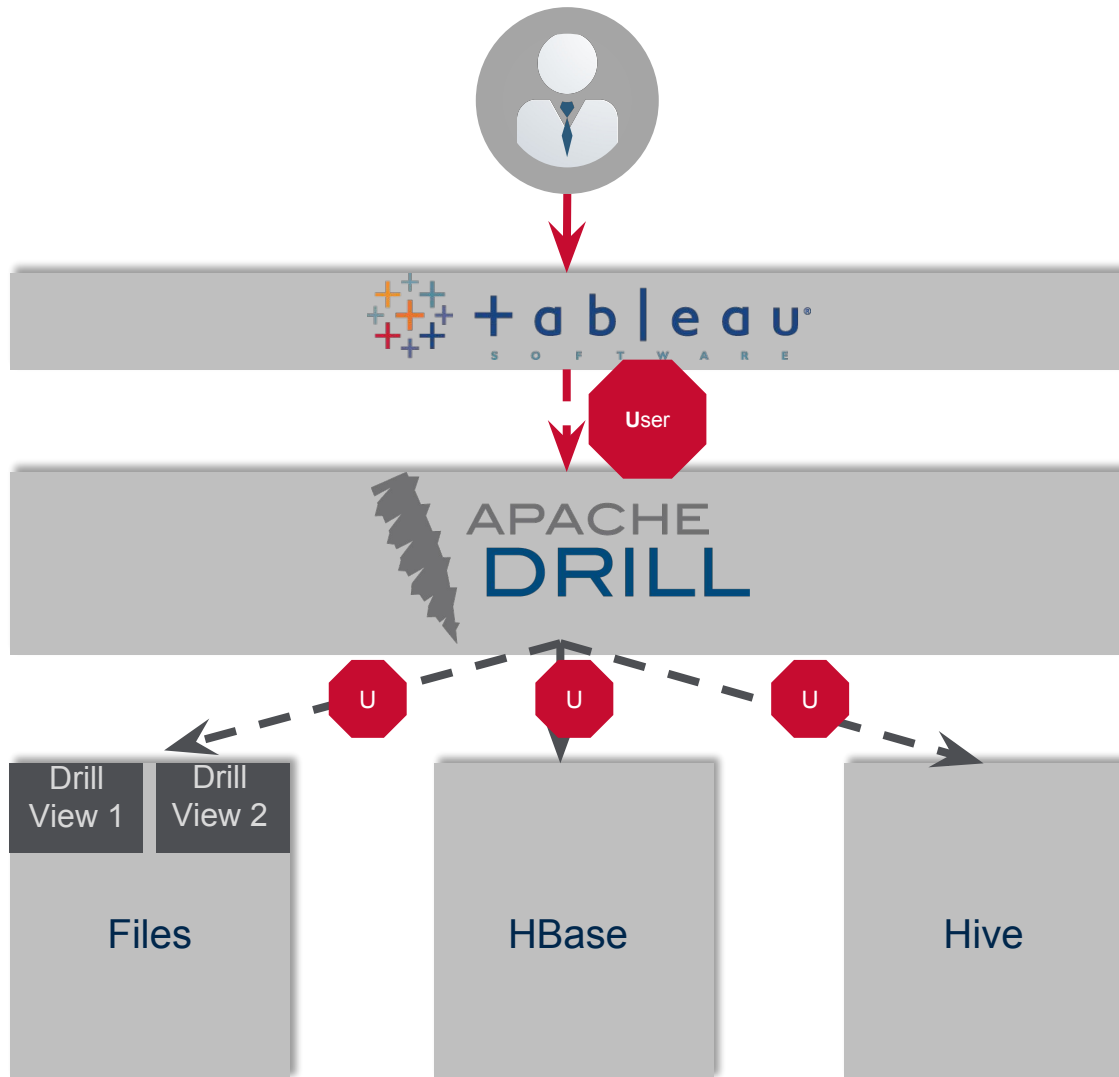


Drill + FS
w/Parquet
(BI/Deep analytics
historical data)



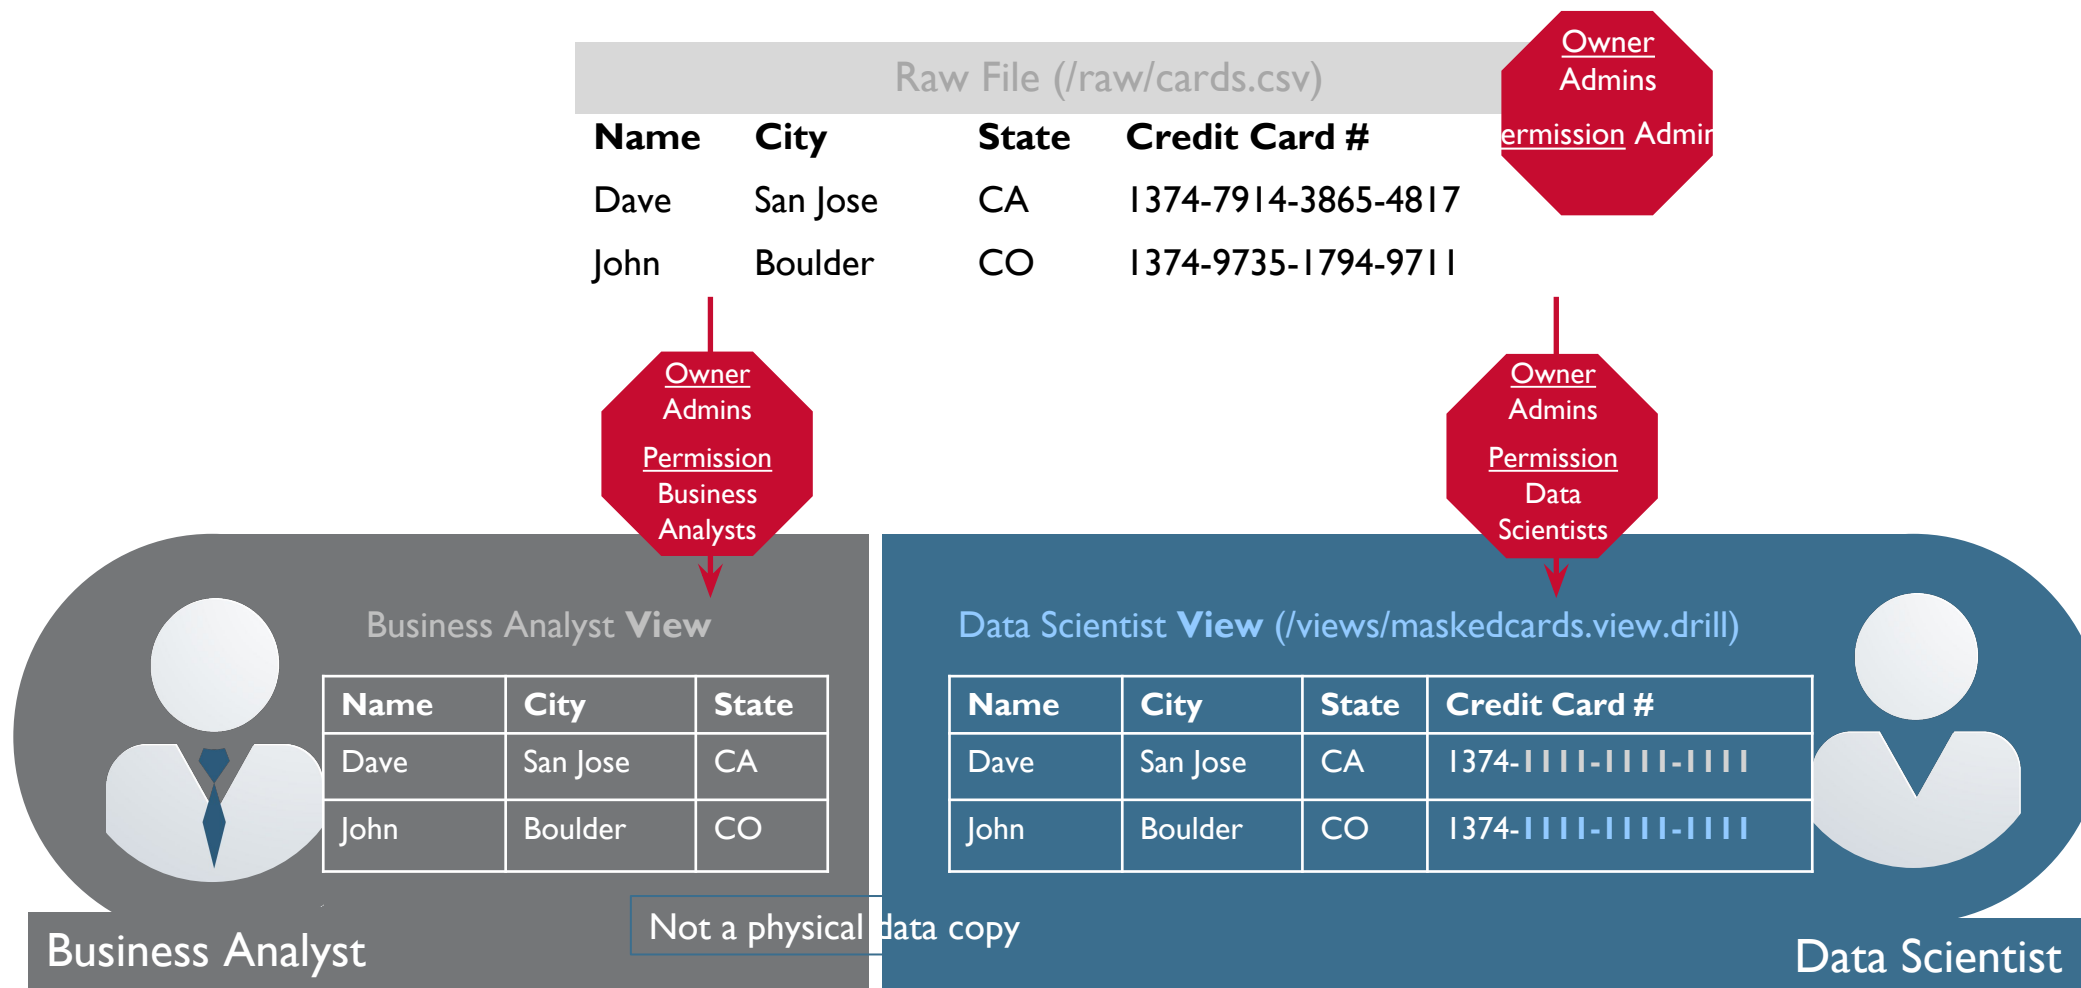
Drill +
MapR-DB/HBase
(Operational
analytics)

End to End security from BI tools to Hadoop

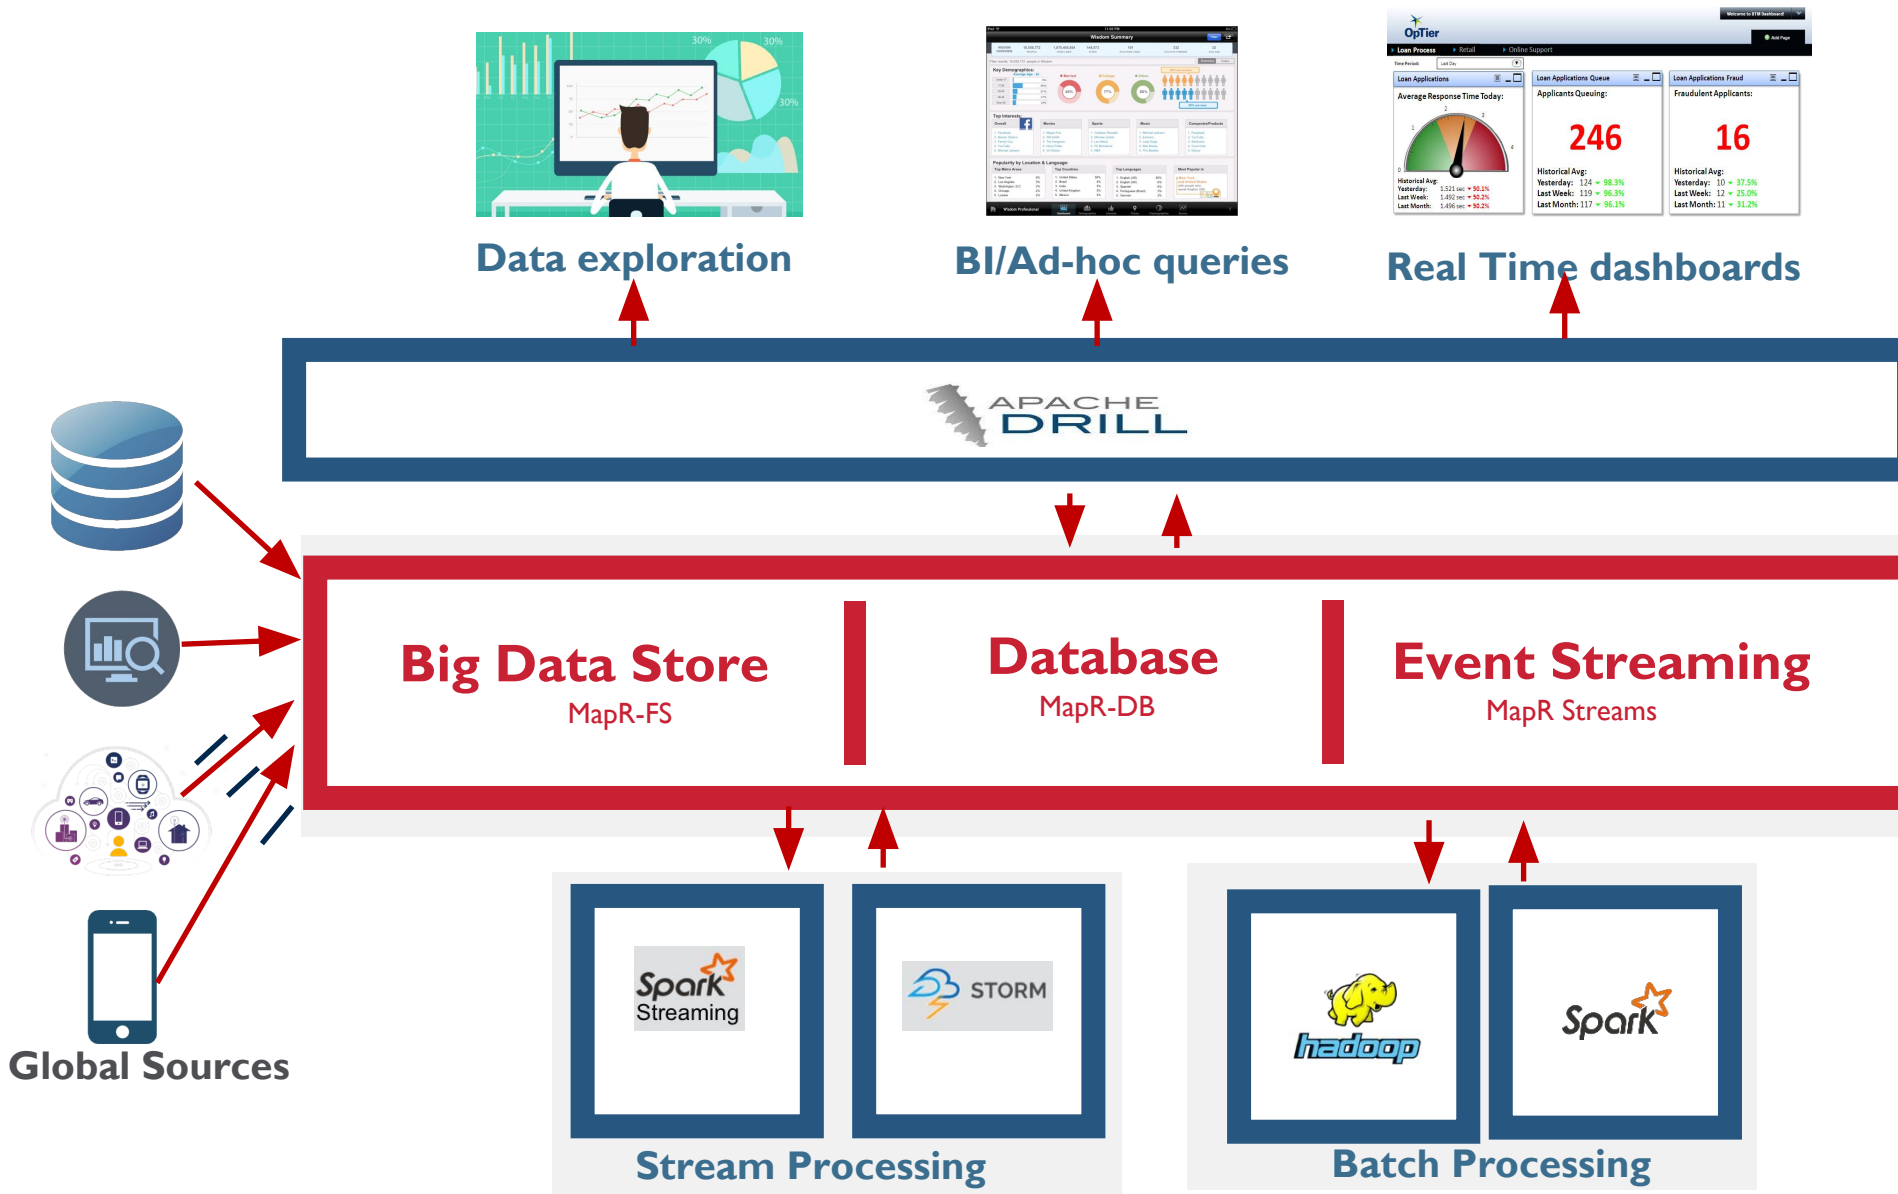


- Security support for JDBC/ODBC drivers & Drill Web UI
- Standard PAM based Authentication
- User impersonation & inbound impersonation
- Fine-grained row and column level access control with Drill Views – no centralized security repository required
- Kerberos & MapR-SASL support between client to Drillbit (Future)
- Encryption support (Future)
 - Between client-Drillbit
 - Between Drillbits

Granular security permissions through Drill views



Evolving towards Unified SQL access layer for MapR platform



- Queries across Files, Tables and Streams
- Real time/Operational analytics
- Schema-less JSON flexibility
- Distributed in-memory SQL engine for high performance at Scale
- Analytics from familiar BI/SQL tools

Drill Best Practices on the MapR Converge Community

Apache Drill Best Practices from the MapR Drill Team

Document created by **Zelaine Fong** on Mar 23, 2016 • Last modified by **Andries Engelbrecht** on Jul 20, 2016



Data Layout and Deployment

Schema Considerations

- ☐ Does Drill have a limit on # of columns?
- ☐ Are there datatypes in Drill that should be favored vs avoided?

Storage Formats

- ☐ What is the preferred storage format for Drill?

Parquet Best Practices

- ☐ What is the preferred compression type when using Parquet with Drill?
- ☐ What is the recommended parquet block size (when running on MapR-FS) for Drill?
- ☐ Can Parquet files created by other tools (e.g., Hive, Spark) be read by Drill?

Data Partitioning

- ☐ How can I partition data in Drill?
- ☐ How do I determine the right partitioning strategy?
- ☐ Are there situations where partition pruning cannot be done in Drill?
- ☐ How do I know if partition pruning has been applied to my Drill query?

Query Tuning and Performance

LIMIT 0 Queries

- ☐ How can I speed up my LIMIT 0 queries on Drill?
- ☐ How do I verify that my Limit 0 Drill queries are benefiting from the optimized code paths?



Hao Zhu
Apr 14, 2016 11:03 AM

★ Correct Answer

For Drill query like the following:

```
select * from table where col in (1,2,3,...);
```

The in-list evaluation is done in a sequential manner. Therefore, if you increase the number of elements in the in-list, the performance of the query could degrade linearly.

If the number of in-list elements contains at least 20 elements, Drill can optimize this query to use an in-memory hash table to store the in-list elements, and perform a table join instead. This optimization can greatly improve the performance.

For example, this query will not benefit from the optimization because the IN list has only 19 elements.

```
select count(1) from t where cast(col as int) in (1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19);
```

Whereas, this one will.

```
select count(1) from t where cast(col as int) in (1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20);
```

To improve the performance of a query with many, but less than 20 in-list elements, you can add duplicate in-list elements so the total number of elements reaches 20.

For example, change the where condition from:

```
"where col in (1,2,3,4,5,6,7,8,9)"
```

To:

```
"where col in (1,2,3,4,5,6,7,8,9,1,1,1,1,1,1,1,1,1,1)"
```

More details are in the following link:

<http://www.openkb.info/2015/08/how-to-improve-performance-of-drill.html>

See the reply in context

<https://community.mapr.com/docs/DOC-1497>

Recommendations On Trying and Using Drill

New to Drill?

- Get started with [Free MapR On Demand training](#)
- [Test Drive Drill](#) on cloud with AWS
- Learn how to use Drill with Hadoop using [MapR sandbox](#)

Ready to play with your data?

- Try out [Apache Drill in 10 mins](#) guide on your desktop
- [Download](#) Drill for your cluster and start exploration
- Comprehensive [tutorials](#) and [documentation](#) available

