

MapR-DB Tablet consistency Check

MapR-DB Tablet consistency Check

This post assumes you have /abizer table created and its empty for us to now use below utility to load the table.

```
1) /opt/mapr/server/tools/loadtest -mode put -numfamilies 1 -numcols 8 -numrows 115450 -table /abizer
23:44:31 0 secs 47391 rows 47391 rows/s 0ms latency 74ms maxLatency
23:44:32 1 secs 69275 rows 21884 rows/s 0ms latency 62ms maxLatency
23:44:33 2 secs 85854 rows 16579 rows/s 1ms latency 332ms maxLatency
23:44:34 3 secs 91451 rows 5597 rows/s 3ms latency 326ms maxLatency
23:44:35 4 secs 99718 rows 8267 rows/s 2ms latency 370ms maxLatency
23:44:36 5 secs 104277 rows 4559 rows/s 4ms latency 665ms maxLatency
23:44:37 6 secs 106115 rows 1838 rows/s 11ms latency 161ms maxLatency
23:44:38 7 secs 115450 rows 9335 rows/s 5ms latency 343ms maxLatency
Overall Rate 14425.84 rows/s, Latency 1ms
```

2) Now check the size and other attributes of table

```
maprcli table info -path /abizer -json
{
  "timestamp":1481356005057,
  "timeofday":"2016-12-09 11:46:45.057 GMT-0800",
  "status":"OK",
  "total":4,
  "data":[
    {
      "path":"/abizer",
      "numregions":4,
      "totallogicalsize":866525184,
      "totalphysicalsize":538484736,
      "totalcopypendingsize":261562368,
      "totalrows":173733,
      "totalnumberofspills":522,
      "totalnumberofsegments":219,
      "autosplit":true,
      "bulkload":false,
      "insertionorder":true,
      "tabletype":"binary",
      "regionsizeemb":4096,
      "audit":false,
      "maxvaluesinmemindex":100,
      "adminaccessperm":"u:root",
      "creatorenamefamilyperm":"u:root",
      "bulkloadperm":"u:root",
      "packperm":"u:root",
      "deletefamilyperm":"u:root",
      "replperm":"u:root",
      "splitmergeperm":"u:root",
      "defaultappendperm":"u:root",
      "defaultcompressionperm":"u:root",
      "defaultmemoryperm":"u:root",
      "defaultreadperm":"u:root",
      "defaultversionperm":"u:root",
      "defaultwriteperm":"u:root",
      "uuid":"315a20f5-ad00-3638-3229-03afb04b5800"
    }
  ]
}
```

3) Now since we know there are 4 regions lets check details of 4 regions (Regions also called Tablet) in the table with respective details.

```
maprcli table region list -path /abizer -json
{
  "timestamp":1481356069293,
  "timeofday":"2016-12-09 11:47:49.293 GMT-0800",
  "status":"OK",
  "total":4,
  "data":[
    {
      "primarymfs":"node10.maprlab.local:5660",
      "secondarymfs": "",
      "startkey":"-INFINITY",
      "endkey":"user3051486366437214287",
      "lastheartbeat":0,
      "fd":"2248.32.131418",
      "logicalsize":215515136,
      "physicalsize":133914624,
      "copypendingsize":127434752,
      "numberofrows":42943,
      "numberofrowswithdelete":0,
    }
  ]
}
```

```

"numberofspills":147,
"numberofsegments":52
},
{
"primarymfs":"node10.maprab.local:5660",
"secondarymfs":"","
"startkey":"user3051486366437214287",
"endkey":"user5105987298073629634",
"lastheartbeat":0,
"fid":"2247.32.131310",
"logicalsize":216645632,
"physicalsize":134545408,
"copypendingsize":0,
"numberofrows":43089,
"numberofrowswithdelete":0,
"numberofspills":149,
"numberofsegments":53
},
{
"primarymfs":"node10.maprab.local:5660",
"secondarymfs":"","
"startkey":"user5105987298073629634",
"endkey":"user720538676648464684",
"lastheartbeat":0,
"fid":"2251.32.131262",
"logicalsize":217849856,
"physicalsize":135897088,
"copypendingsize":0,
"numberofrows":44198,
"numberofrowswithdelete":0,
"numberofspills":112,
"numberofsegments":56
},
{
"primarymfs":"node10.maprab.local:5660",
"secondarymfs":"","
"startkey":"user720538676648464684",
"endkey":"INFINITY",
"lastheartbeat":0,
"fid":"2250.32.131186",
"logicalsize":216514560,
"physicalsize":134127616,
"copypendingsize":134127616,
"numberofrows":43503,
"numberofrowswithdelete":0,
"numberofspills":114,
"numberofsegments":58
}
]
}

```

4) To get Info/details of local tablets per node alternatively can also be viewed by mrconfig commands.

```
/opt/mapr/server/mrconfig dbinfo tablets
```

```
-----
|From Instance 5660:|
-----
```

```

tablet 2250.32.131186 nref 0 npartitions 1 logicalMB 206 physicalMB 127 rows 43503 splitState None attrAutoSplit 1 tabletSplitThreshSizeMB 6144 partitionSplit1
updateError 0
tablet 2251.32.131262 nref 0 npartitions 1 logicalMB 207 physicalMB 129 rows 44198 splitState None attrAutoSplit 1 tabletSplitThreshSizeMB 6144 partitionSplit1
updateError 0
tablet 2247.32.131310 nref 0 npartitions 2 logicalMB 206 physicalMB 128 rows 43089 splitState None attrAutoSplit 1 tabletSplitThreshSizeMB 6144 partitionSplit1
updateError 0
tablet 2248.32.131418 nref 0 npartitions 2 logicalMB 205 physicalMB 127 rows 42943 splitState None attrAutoSplit 1 tabletSplitThreshSizeMB 6144 partitionSplit1
updateError 0

```

5) Lets say we believe "2251.32.131262" fid has some issues and we want to review the stats of the fid in question, below command can give the info needed.

```
[root@node10 ~]# maprccli debugdb statTablet -fid 2251.32.131262 -json
```

```

{
"timestamp":1481356575116,
"timeofday":"2016-12-09 11:56:15.116 GMT-0800",
"status":"OK",
"total":1,
"data":1,
{
{
"numPhysicalBlocks":16589,
"numLogicalBlocks":26593,
"numRows":44198,
"numRowsWithDelete":0,
"numRemoteBlocks":0,
"numSpills":112,
"numSegments":56
}
}
}

```

```
]
}
```

6) To check the consistency of the tablet we can run below command which check and reports details related to tablet. The startkey and endkey needs to be from ou /abizer -json" for the region/Tablet in question.

If this fails for any reason we need to check /opt/mapr/logs/mfs.log-5 on the master node for the tablet and see if we see the error reported to debug further but norm see clean output as below.

```
[root@node10 ~]# maprcli debugdb checkTablet -fid 2251.32.131262 -tracefile /tmp/tablet-2251.32.131262 -startkey 'user5105987298073629634' -endkey 'user7204'
```

TabletRangeCheck done

```
[root@node10 ~]# cat /tmp/tablet-2251.32.131262
Tablet 2251.32.131262
Segmap 2251.607.132412
spill 2251.266.131730 parent 2251.265.131728 idx 0 size 2048000
keymap 2251.266.131730 parent 2251.265.131728 idx 0 offset 12 len 31471
cf 2251.266.131730 parent 2251.265.131728 idx 0 family 1 offset 524288 len 2110718
cf 2251.266.131730 parent 2251.265.131728 idx 0 family 2 offset 2686976 len 1048712
Spillmap 2251.265.131728 parent 2251.607.132412
spill 2251.390.131978 parent 2251.383.131964 idx 0 size 2088960
keymap 2251.390.131978 parent 2251.383.131964 idx 0 offset 12 len 32489
cf 2251.390.131978 parent 2251.383.131964 idx 0 family 1 offset 524288 len 2185207
cf 2251.390.131978 parent 2251.383.131964 idx 0 family 2 offset 2752512 len 1050456
Spillmap 2251.383.131964 parent 2251.607.132412
spill 2251.546.132290 parent 2251.383.131964 idx 1 size 1474560
keymap 2251.546.132290 parent 2251.383.131964 idx 1 offset 12 len 19507
cf 2251.546.132290 parent 2251.383.131964 idx 1 family 1 offset 524288 len 1138369
cf 2251.546.132290 parent 2251.383.131964 idx 1 family 2 offset 1703936 len 1138798
[root@node10 ~]#
```

Above file proves everything is fine with the tablet and there are no inconsistencies found.

Posted by Abizer Adenwala at 9:47 AM

