

2nd April 2015 The Number of Maximum Attempts of an Yarn Application in Hadoop Two

A hadoop job in Yarn is called an application. For example, a mapreduce job is one-to-one mapped to an application currently. If an application fails for a certain reason, Yarn will retry the application, i.e., launch a new application master to manage the lifecycle of the application. The number of maximum attempts is set by configuration "*yarn.resourcemanager.am.max-attempts*" with a default value **2** and it's a global setting for all application masters.

```
public static final String RM_AM_MAX_ATTEMPTS =
    RM_PREFIX + "am.max-attempts";
public static final int DEFAULT_RM_AM_MAX_ATTEMPTS = 2;
```

However, each application can set its own maximum number as long as the individual number is less than the global upper bound. Otherwise, the resourcemanager will override it. The `ApplicationSubmissionContext` class provides a method to set this number.

```
public abstract class ApplicationSubmissionContext {

    /**
     * Set the number of max attempts of the application to be submitted. WARNING:
     * it should be no larger than the global number of max attempts in the Yarn
     * configuration.
     * @param maxAppAttempts the number of max attempts of the application
     * to be submitted.
     */
    @Public
    @Stable
    public abstract void setMaxAppAttempts(int maxAppAttempts);

}
```

The retry logic lives inside the class `RMAppImpl`, which is a representation of an Yarn application on the resource manager side.

```
public class RMAppImpl implements RMApp, Recoverable {
    private final int maxAppAttempts;
    private boolean isNumAttemptsBeyondThreshold = false;

    public RMAppImpl(ApplicationId applicationId, RMContext rmContext,
        Configuration config, String name, String user, String queue,
        ApplicationSubmissionContext submissionContext, YarnScheduler scheduler,
        ApplicationMasterService masterService, long submitTime,
        String applicationType, Set applicationTags,
        ResourceRequest amReq) {

        int globalMaxAppAttempts = conf.getInt(YarnConfiguration.RM_AM_MAX_ATTEMPTS,
            YarnConfiguration.DEFAULT_RM_AM_MAX_ATTEMPTS);
        int individualMaxAppAttempts = submissionContext.getMaxAppAttempts();
        if (individualMaxAppAttempts <= 0 ||
            individualMaxAppAttempts > globalMaxAppAttempts) {
            this.maxAppAttempts = globalMaxAppAttempts;
            LOG.warn("The specific max attempts: " + individualMaxAppAttempts
```

```

        + " for application: " + applicationId.getId()
        + " is invalid, because it is out of the range [1, "
        + globalMaxAppAttempts + "]. Use the global max attempts instead.");
    } else {
        this.maxAppAttempts = individualMaxAppAttempts;
    }
    ...
}
}

```

As we can see, the *globalMaxAppAttempts* is obtained from Yarn configuration file and the *maxAppAttempts* is obtained from *ApplicationSubmissionContext* and then compared with the *globalMaxAppAttempts*.

The number of actual failed application attempts is calculated by the following method to exclude AM preempt hardware failures or NM resync.

```

private int getNumFailedAppAttempts() {
    int completedAttempts = 0;
    long endTime = this.systemClock.getTime();
    // Do not count AM preemption, hardware failures or NM resync
    // as attempt failure.
    for (RMAAppAttempt attempt : attempts.values()) {
        if (attempt.shouldCountTowardsMaxAttemptRetry()) {
            if (this.attemptFailuresValidityInterval <= 0
                || (attempt.getFinishTime() > endTime
                    - this.attemptFailuresValidityInterval)) {
                completedAttempts++;
            }
        }
    }
    return completedAttempts;
}

```

where *shouldCountTowardsMaxAttemptRetry()* is defined in class *RMAAppAttemptImpl* as follows.

```

public boolean shouldCountTowardsMaxAttemptRetry() {
    try {
        this.readLock.lock();
        int exitStatus = getAMContainerExitStatus();
        return !(exitStatus == ContainerExitStatus.PREEMPTED
            || exitStatus == ContainerExitStatus.ABORTED
            || exitStatus == ContainerExitStatus.DISKS_FAILED
            || exitStatus == ContainerExitStatus.KILLED_BY_RESOURCEMANAGER);
    } finally {
        this.readLock.unlock();
    }
}

```

Then the variable *maxAppAttempts* is checked in the state transition *AttemptFailedTransition* to set the *isNumAttemptsBeyondThreshold*.

```

private static final class AttemptFailedTransition implements
    MultipleArcTransition {

    private final RMapAppState initialState;

    @Override
    public RMapAppState transition(RMapAppImpl app, RMapAppEvent event) {
        int numberOfFailure = app.getNumFailedAppAttempts();
        LOG.info("The number of failed attempts"
            + (app.attemptFailuresValidityInterval > 0 ? " in previous "
                + app.attemptFailuresValidityInterval + " milliseconds " : " ")
            + "is " + numberOfFailure + ". The max attempts is "
            + app.maxAppAttempts);
        if (!app.submissionContext.getUnmanagedAM()
            && numberOfFailure < app.maxAppAttempts) {
            if (initialState.equals(RMapAppState.KILLING)) {
                // If this is not last attempt, app should be killed instead of
                // launching a new attempt
                app.rememberTargetTransitionsAndStoreState(event,
                    new AppKilledTransition(), RMapAppState.KILLED, RMapAppState.KILLED);
                return RMapAppState.FINAL_SAVING;
            }

            boolean transferStateFromPreviousAttempt;
            RMapAppFailedAttemptEvent failedEvent = (RMapAppFailedAttemptEvent) event;
            transferStateFromPreviousAttempt =
                failedEvent.getTransferStateFromPreviousAttempt();

            RMapAppAttempt oldAttempt = app.currentAttempt;
            app.createAndStartNewAttempt(transferStateFromPreviousAttempt);
            // Transfer the state from the previous attempt to the current attempt.
            // Note that the previous failed attempt may still be collecting the
            // container events from the scheduler and update its data structures
            // before the new attempt is created. We always transferState for
            // finished containers so that they can be acked to NM,
            // but when pulling finished container we will check this flag again.
            ((RMapAppAttemptImpl) app.currentAttempt)
                .transferStateFromPreviousAttempt(oldAttempt);
            return initialState;
        } else {
            if (numberOfFailure >= app.maxAppAttempts) {
                app.isNumAttemptsBeyondThreshold = true;
            }
            app.rememberTargetTransitionsAndStoreState(event,
                new AttemptFailedFinalStateSavedTransition(), RMapAppState.FAILED,
                RMapAppState.FAILED);
            return RMapAppState.FINAL_SAVING;
        }
    }
}

```

As shown above, if *numberOfFailure* >= *maxAppAttempts*, the application state transfers to **FAILED** and no more retri