

## Index

- [Creating the Kubernetes cluster](#)
  - [Getting the Kops tool](#)
  - [Creating a Key Pair](#)
  - [Creating the Kube cluster](#)
    - [Create the cluster configs](#)
    - [Get cluster](#)
  - [Edit the Cluster Configs](#)
  - [Get the current instance groups](#)
    - [Build the Cluster](#)
- [kubectl](#)
  - [Getting kubectl tool](#)
  - [Connection to the Kube API with the Kubectl tool](#)
- [General Kubernetes Usage](#)
  - [Listing Pods](#)
- [You have a Kubernetes cluster, now what?](#)
  - [Launching Ingress](#)
- [Build and Deploy With Wercker](#)
  - [Triggering your first build](#)
  - [Configuring continuous deployments](#)
  - [Conclusion](#)

## Creating the Kubernetes cluster

### Getting the Kops tool

The list of the latest downloads for kops can be found here:

<https://github.com/kubernetes/kops/releases/>

```
$ chmod +x kops-linux-amd64
$ mv kops-linux-amd64 /usr/local/bin/
```

Confirm that it is properly installed by typing kops. You should get a similar output:

```
[kopsuser@kops-demo ~]$ kops
kops is kubernetes ops.
It allows you to create, destroy, upgrade and maintain clusters.

Usage:
  kops [command]

Available Commands:
  create      create resources
  delete      delete clusters
  describe    describe objects
  edit        edit items
  export      export clusters/kubecfg
  get         list or get objects
  import      import clusters
  rolling-update rolling update clusters
  secrets      Manage secrets & keys
  toolbox     Misc infrequently used commands
  update      update clusters
  upgrade      upgrade clusters
  version     Print the client version information

Flags:
  --alsologtostderr      log to standard error as well as files
  --config string         config file (default is $HOME/.kops.yaml)
  -h, --help              help for kops
  --log_backtrace_at traceLocation when logging hits line file:N, emit a stack trace (default :0)
  --log_dir string        If non-empty, write log files in this directory
  --logtostderr           log to standard error instead of files (default false)
  --name string           Name of cluster
  --state string          Location of state storage
  --stderrthreshold severity logs at or above this threshold go to stderr (default 2)
  -v, --v Level           log level for V logs
  --vmodule moduleSpec    comma-separated list of pattern=N settings for file-filtered logging

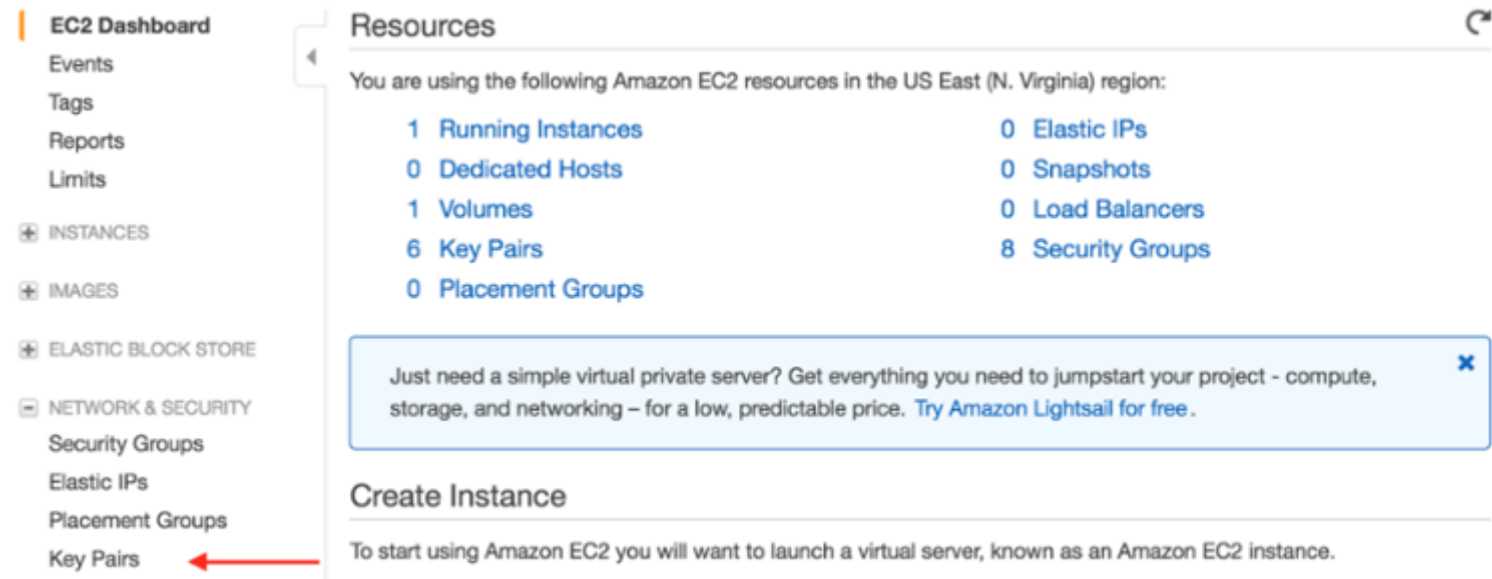
Use "kops [command] --help" for more information about a command.
[kopsuser@kops-demo ~]$
```

## Creating a Key Pair

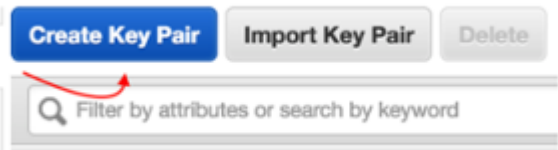
**Please note:** you will need to create your own ssh public key and import it to the Key Pairs section under EC2 in the Amazon Web Console.

Before we create our kube cluster we will need to create our SSH key pair needed for access:

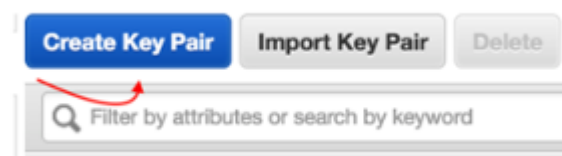
- 1. Log into your [AWS web console](#).
- 2. Select the EC2 option.
- 3. Under the Network & Security option select Key Pairs.



- 4. Select Create Key Pair.



- 5. Give your Key Pair a name and click create.



The private key file is automatically downloaded by your browser. The filename is the name you specified as the name of your key pair, and the filename extension is .pem. **Save the private key file in a safe place!**

**Important:** This is the only chance for you to save the private key file. You'll need to provide the name of your key pair when you launch an instance and the corresponding private key each time you connect to the instance.

6. If you use an SSH client on a Mac or Linux computer to connect to your Linux instance, use the following command to set the permissions of your private key file so that only you can read it:

```
$ chmod 400 my-key-pair.pem
```

## Creating the Kube cluster

Please Note: Before we can create our kube cluster we will have to:

- Create a route53 domain for our cluster.
- Create an S3 bucket to store our clusters state with the following name: kop-store-**<your\_name>**.

Instructions for accomplishing both can be found at steps 2 and 3 at the link below:

<https://kubernetes.io/docs/getting-started-guides/kops/>

Export the various env variables:

```
export AWS_ACCESS_KEY_ID=""
export AWS_SECRET_ACCESS_KEY=""
export AWS_DEFAULT_REGION=us-east-1
export KOPS_STATE_STORE=s3://kop-store-<your_name>
export NAME=demo.k8s.<your_valid_domain_name>
Example of a valid domain name is: demo.k8s.wercker.com
```

Create the cluster configs:

```
kops create cluster
--name=${NAME}
--cloud=aws
--kubernetes-version=1.5.1
--zones=us-east-1b,us-east-1c,us-east-1d
--ssh-public-key=my-key-pair.pem
--network-cidr=172.30.0.0/16
--admin-access=172.30.0.0/16
--node-count=1
--master-zones=us-east-1b,us-east-1c,us-east-1d
--master-size=t2.medium
```

Get cluster

The below command will show you your newly created cluster:

```
$ kops get cluster
NAME CLOUD ZONES
Demo.k8s.wercker.com aws us-east-1b
```

Edit the Cluster Configs:

```
kops edit cluster ${NAME}
```

Get the current instance groups:

```
$ kops get instancegroups --name ${NAME}
```

```
[kopsuser@kops-demo ~]$ kops get instancegroups --name ${NAME}
NAME          ROLE    MACHINETYPE  MIN  MAX  ZONES
master-us-east-1b  Master  t2.medium    1    1    us-east-1b
master-us-east-1c  Master  t2.medium    1    1    us-east-1c
master-us-east-1d  Master  t2.medium    1    1    us-east-1d
nodes          Node    t2.medium    1    1    us-east-1b,us-east-1c,us-east-1d
[kopsuser@kops-demo ~]$
```

This output shows you that it will create 4 nodes all together. There will be 3 master nodes for redundancy and 1 minion node.

Build the Cluster

This will start launching: subnets, ec2 instances, and start creating your cluster:

```
kops update cluster ${NAME} --yes
```

```
[kopsuser@kops-demo ~]$ kops update cluster ${NAME} --yes
I0113 01:21:54.349096 56789 populate_cluster_spec.go:196] Defaulting DNS zone to: ZGRP7R89GCXW2
I0113 01:21:57.619006 56789 executor.go:68] Tasks: 0 done / 63 total; 30 can run
I0113 01:21:58.045503 56789 vfs_castore.go:384] Issuing new certificate: "kubelet"
I0113 01:21:58.702519 56789 vfs_castore.go:384] Issuing new certificate: "master"
I0113 01:21:58.790586 56789 vfs_castore.go:384] Issuing new certificate: "kubecfg"
I0113 01:21:59.358875 56789 executor.go:68] Tasks: 30 done / 63 total; 12 can run
I0113 01:22:00.433470 56789 executor.go:68] Tasks: 42 done / 63 total; 17 can run
I0113 01:22:01.729195 56789 launchconfiguration.go:276] Waiting for IAM to replicate
I0113 01:22:01.748502 56789 launchconfiguration.go:276] Waiting for IAM to replicate
I0113 01:22:02.060601 56789 launchconfiguration.go:276] Waiting for IAM to replicate
I0113 01:22:02.075549 56789 launchconfiguration.go:276] Waiting for IAM to replicate
I0113 01:22:12.736906 56789 executor.go:68] Tasks: 59 done / 63 total; 4 can run
I0113 01:22:13.435753 56789 executor.go:68] Tasks: 63 done / 63 total; 0 can run
I0113 01:22:13.646923 56789 update_cluster.go:150] Exporting kubecfg for cluster
Wrote config for demo.k8s.flowlog-stats.com to "/home/kopsuser/.kube/config"

Cluster is starting. It should be ready in a few minutes.

Suggestions:
* list nodes: kubectl get nodes --show-labels
* ssh to the master: ssh -i ~/.ssh/id_rsa admin@api.demo.k8s.flowlog-stats.com
* read about installing addons: https://github.com/kubernetes/kops/blob/master/docs/addons.md
```

Once you build the cluster the kube config file with the authorization credentials for access will be in your home directory:

```
~/.kube/config
```

```
[kopsuser@kops-demo ~]$ ls -a
.  ..  .bash_logout  .bash_profile  .bashrc  .kube  .pki  .ssh
[kopsuser@kops-demo ~]$
```

```
[kopsuser@kops-demo .kube]$ pwd && ls
/home/kopsuser/.kube
config
[kopsuser@kops-demo .kube]$
```

# kubect

## Getting kubectl tool

Download the kubectl command line tool:

```
$ curl -LO https://storage.googleapis.com/kubernetes-release/release/$(curl -s https://storage.g
```

Make the kubectl binary executable and move it to your PATH:

```
$ chmod +x kubectl
$ sudo mv ./kubectl /usr/local/bin/kubectl
```

After downloading let's enable shell autocompletion:

```
$ echo "source <(kubectl completion bash)" >> ~/.bashrc
```

Once you've moved the kubectl binary to `/usr/local/bin/` confirm that it is properly installed by typing the below command. You should get a similar output from the screen capture below:

```
$ kubectl
```

```
kops@server:kops-demo ~$ kubectl
kubectl controls the Kubernetes cluster manager.

Find more information at https://github.com/kubernetes/kubernetes.

Basic Commands (Beginner):
  create      Create a resource by filename or stdin
  expose      Take a replication controller, service, deployment or pod and expose it as a new Kubernetes Service
  run         Run a particular image on the cluster
  set         Set specific features on objects

Basic Commands (Intermediate):
  get         Display one or many resources
  explain     Documentation of resources
  edit        Edit a resource on the server
  delete      Delete resources by filenames, stdin, resources and names, or by resources and label selector

Deploy Commands:
  rollout     Manage a deployment rollout
  rolling-update Perform a rolling update of the given ReplicationController
  scale       Set a new size for a Deployment, ReplicaSet, Replication Controller, or Job
  autoscale   Auto-scale a Deployment, ReplicaSet, or ReplicationController

Cluster Management Commands:
  certificate Modify certificate resources.
  cluster-info Display cluster info
  top         Display Resource (CPU/Memory/Storage) usage
  cordon      Mark node as unschedulable
  uncordon    Mark node as schedulable
  drain       Drain node in preparation for maintenance
  taint       Update the taints on one or more nodes

Troubleshooting and Debugging Commands:
  describe    Show details of a specific resource or group of resources
  logs        Print the logs for a container in a pod
  attach      Attach to a running container
  exec        Execute a command in a container
  port-forward Forward one or more local ports to a pod
  proxy       Run a proxy to the Kubernetes API server
  cp          Copy files and directories to and from containers.

Advanced Commands:
  apply       Apply a configuration to a resource by filename or stdin
  patch       Update field(s) of a resource using strategic merge patch
  replace     Replace a resource by filename or stdin
```

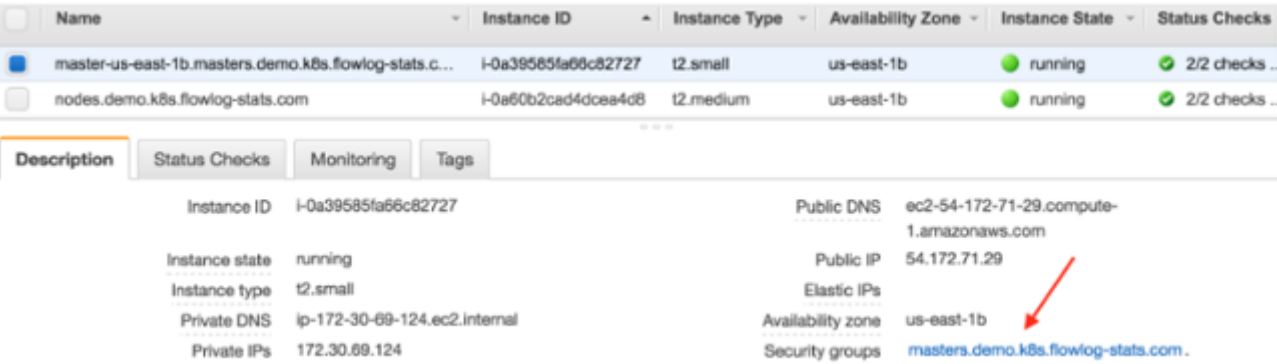
## Connection to the Kube API with the Kubectl tool

The kubectl tool uses a kube config file (default location is in `~/.kube/config` ) to connect and to communicate with our kube cluster. The file that Kops generates is pointing to the public IP address of the cluster.

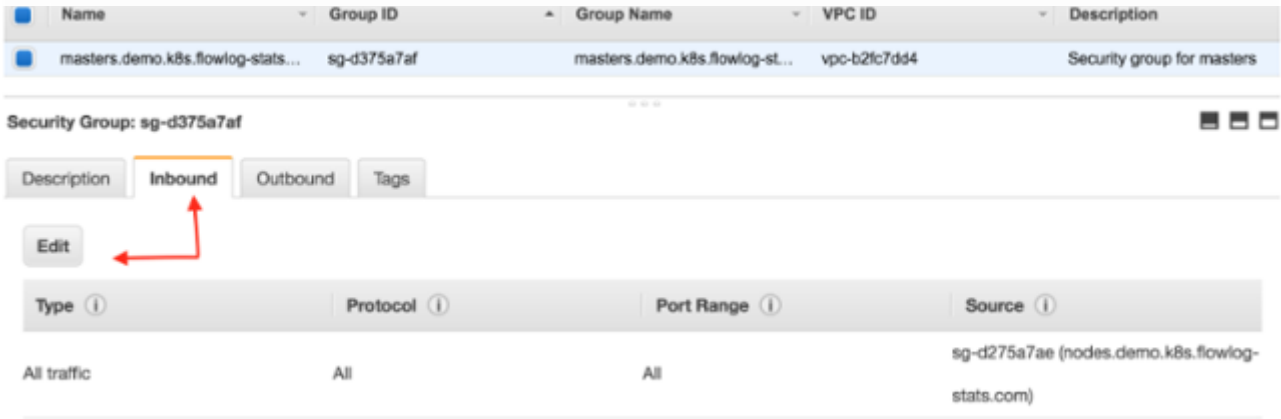
We have set the security groups within AWS so that you cannot reach the public IP of the kube API. With that, you have to open up access to the kube API. **We want to severely limit access to these servers.**

At this point we will need to add our public IP address to the inbound rule for our master instance in the Security Group of our AWS console.

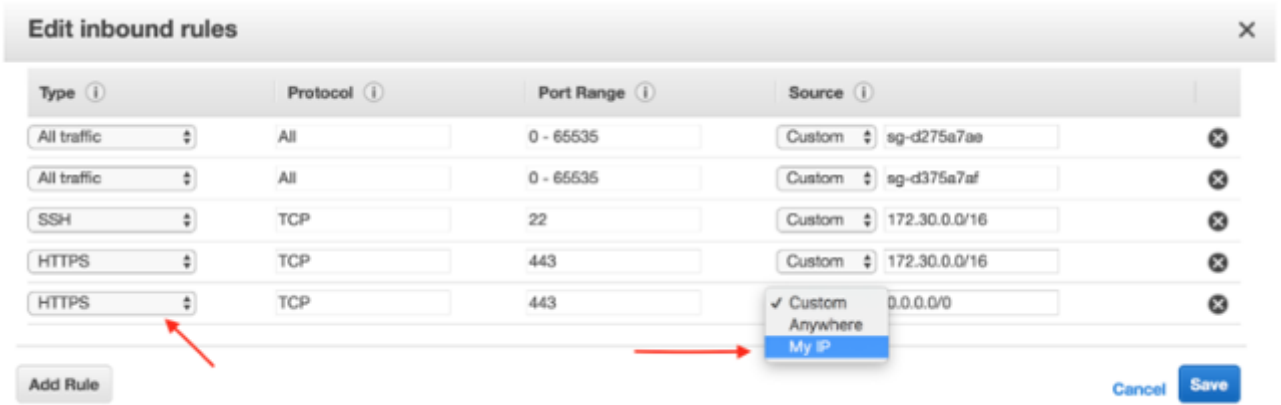
- Log into your [AWS Console](#) and select EC2.
- Select the master instance and click on its Security Group option.



- Select the Inbound tab then Edit.



- Choose HTTPS for the type and select My IP then save.



- Before we leave the Security Group section let us repeat the steps above and [whitelist the below IP addresses](#). We will need these IP's for one of our steps moving forward that allows the Wercker.com CI/CD service to reach our Kubernetes cluster.

After you've whitelisted your IP verify connectivity by checking with the nslookup command:

```
$ nslookup api.demo.k8s.wercker.com
```

Please Note: This can take up to 15 minutes for the DNS records to update with a response.

```
[kopsuser@kops-demo ~]$ nslookup api.demo.k8s.flowlog-stats.com
Server:      10.0.1.1
Address:     10.0.1.1#53

Non-authoritative answer:
Name:   api.demo.k8s.flowlog-stats.com
Address: 52.87.252.76
```

# General Kubernetes Usage

## Listing Pods

You can tell the command line which namespace you want to look at.

Another interesting namespace is the **kube-system** namespace. This is a namespace that is automatically created for you and it runs Kubernetes system pods in it.

```
kubectl --namespace kube-system get pods -o wide
```



NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
dns-controller-1331203010-lhqs9	1/1	Running	0	25m	10.123.45.2	ip-172-30-69-124.ec2.internal
etcd-server-events-ip-172-30-69-124.ec2.internal	1/1	Running	4	26m	172.30.69.124	ip-172-30-69-124.ec2.internal
etcd-server-ip-172-30-69-124.ec2.internal	1/1	Running	4	26m	172.30.69.124	ip-172-30-69-124.ec2.internal
kube-apiserver-ip-172-30-69-124.ec2.internal	1/1	Running	5	26m	172.30.69.124	ip-172-30-69-124.ec2.internal
kube-controller-manager-ip-172-30-69-124.ec2.internal	1/1	Running	0	26m	172.30.69.124	ip-172-30-69-124.ec2.internal
kube-dns-v20-3531996453-kcw6	3/3	Running	0	25m	100.96.1.2	ip-172-30-64-169.ec2.internal
kube-dns-v20-3531996453-m3l4p	3/3	Running	0	25m	100.96.1.3	ip-172-30-64-169.ec2.internal
kube-proxy-ip-172-30-64-169.ec2.internal	1/1	Running	0	24m	172.30.64.169	ip-172-30-64-169.ec2.internal
kube-proxy-ip-172-30-69-124.ec2.internal	1/1	Running	0	26m	172.30.69.124	ip-172-30-69-124.ec2.internal
kube-scheduler-ip-172-30-69-124.ec2.internal	1/1	Running	0	26m	172.30.69.124	ip-172-30-69-124.ec2.internal

Also notice there is a -o wide at the end of the command. That is optional but it gives extra information on which node the pod is running on.

# You have a Kubernetes cluster, now what?

We now need to launch some stuff into the cluster and we will use the below repo as an example:

<https://github.com/wercker/kubernetes-ci-cd>

First clone the folder to your system:

```
$ git clone https://github.com/wercker/kubernetes-ci-cd
$ cd kubernetes-ci-cd
```

Create the files needed with the kubectl create -f command for each of the yaml files below:

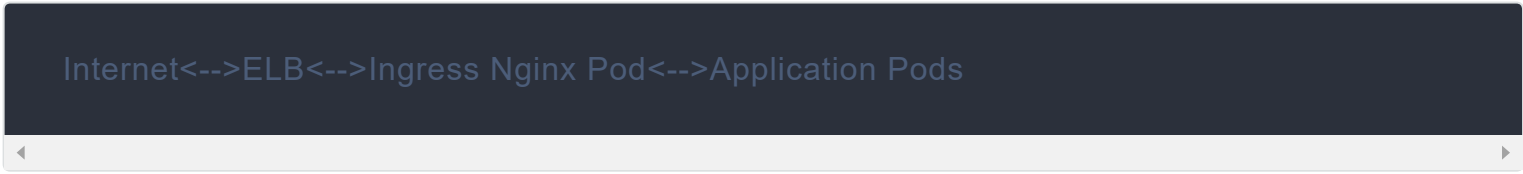
```
$ kubectl --namespace default create -f ./kubernetes/ingress/
```

This will create all of the resources in the `ingress` directory.

## Launching Ingress

The Kubernetes Ingress is a Nginx pod that reads from the Kubernetes API to create it's own configuration so that it can dynamically route traffic inbound to the cluster.

The topology is as follows:



The ingress is created after the creation of the default-backend.yaml and rc.yaml files.

Verify the creation of the ingress (nginx) with:

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
ingress-controller-2171626210-2bg8c  1/1     Running   0          4m
ingress-default-backend-2249275363-8f0tw  1/1     Running   0          4m
```

# Build and Deploy With Wercker

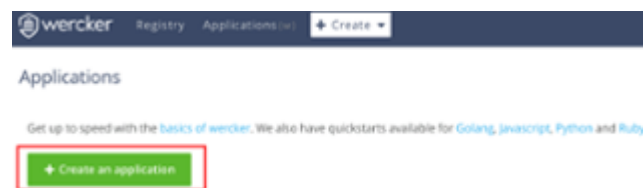
You can follow this step-by-step guide on developing, building and deploying a sample app with Wercker:

- <http://devcenter.wercker.com/docs/quickstarts/building/golang>

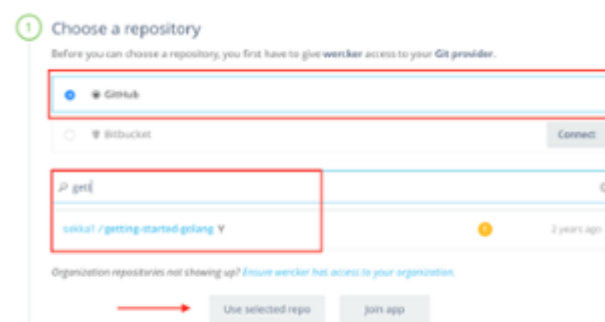


For our purposes we will walk through adding your app to wercker:

- [Sign up for a wercker.com account](#) (you can also sign in with your GitHub account)
- After signing in select the option to create a new Application.



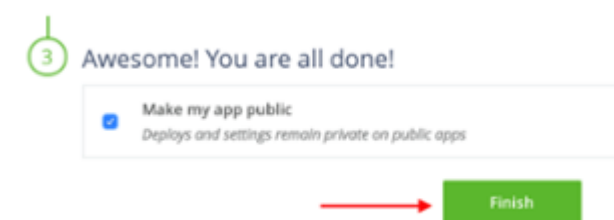
- Select your Git provider along with the App we will deploy and lastly use selected repo.



- Choose Wercker will checkout the code without using an SSH key.



- Check the box and select Finish!

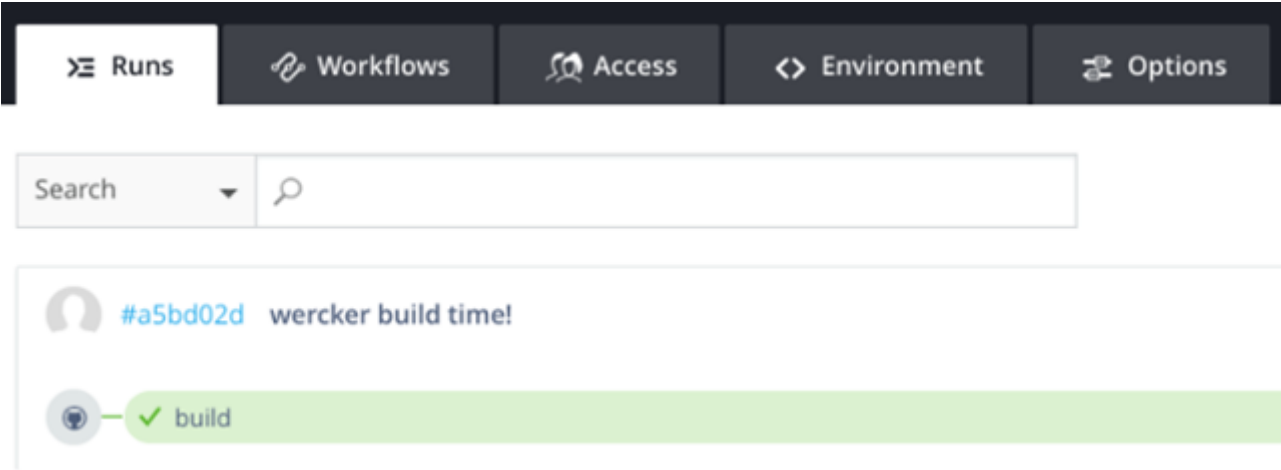


## Triggering your first build

- Make a change to one of the files in the getting-started-golang folder from the command line then run the below commands:

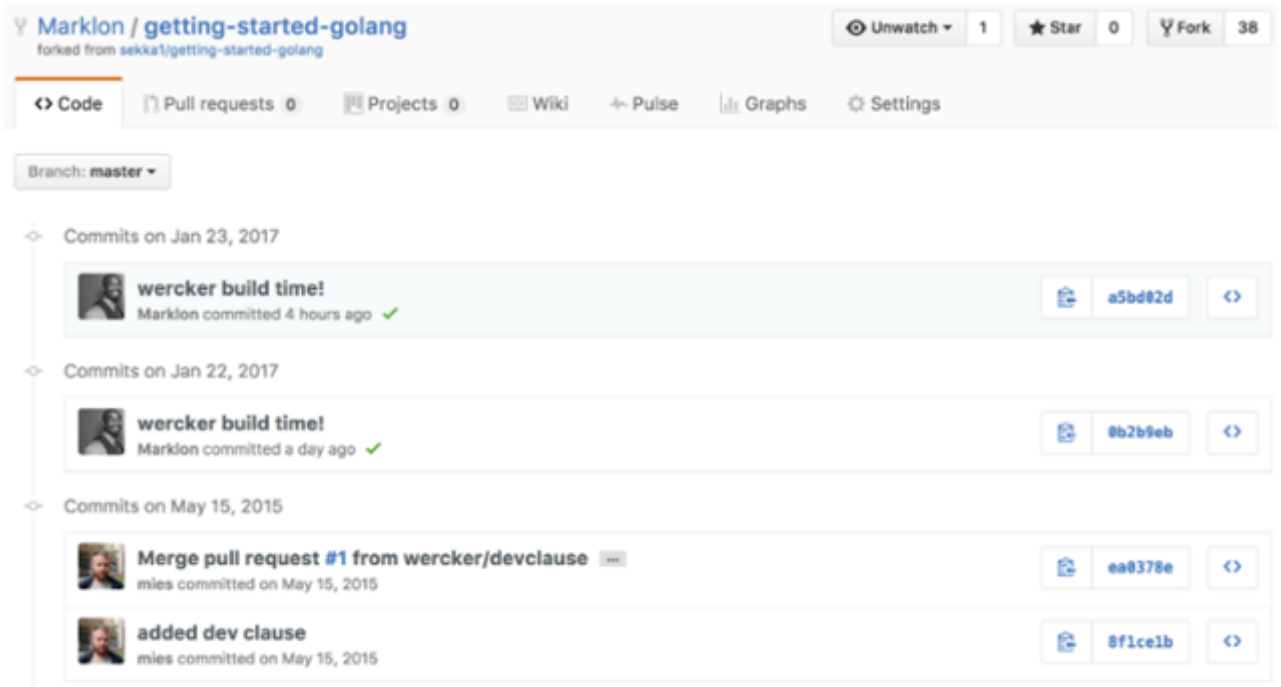
```
$ git commit -am 'wercker build time!'
$ git push origin master
```

Finally, navigate to your app page where you will see a new build has been triggered!



You can also see the commits from wercker in the GitHub repo:

<https://github.com/Marklon/getting-started-golang>



## Configuring continuous deployments

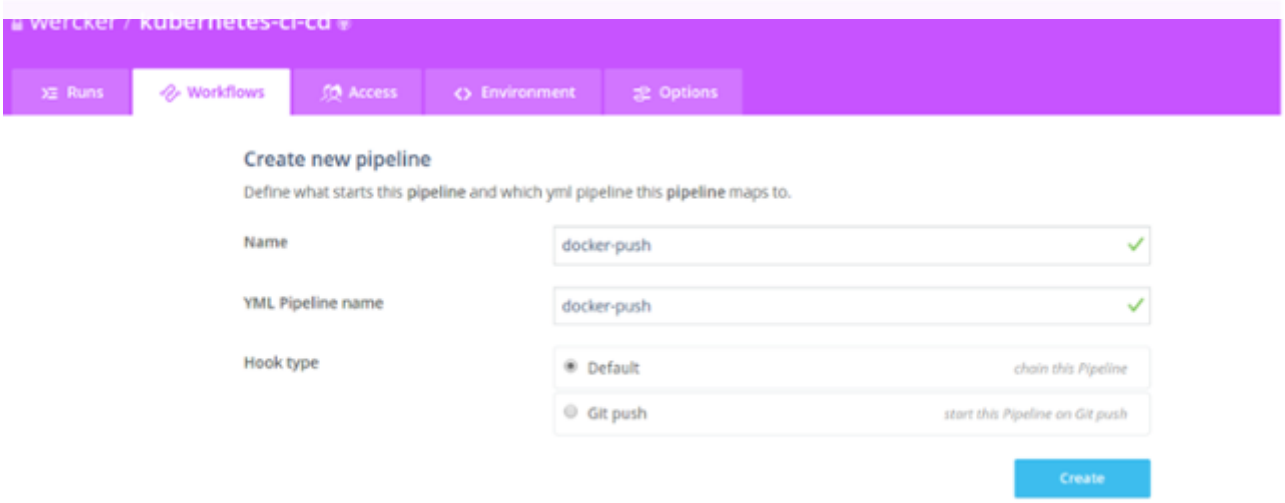
We will add 2 actions into Wercker.com that is in our wercker.yml file.

1) Push the built docker container to our [quay.io repository](https://quay.io). You can also push it to any other Docker repositories.

a) Our definition for the “push docker” action: <https://github.com/wercker/kubernetes-ci-cd/blob/master/wercker.yml#L42>.

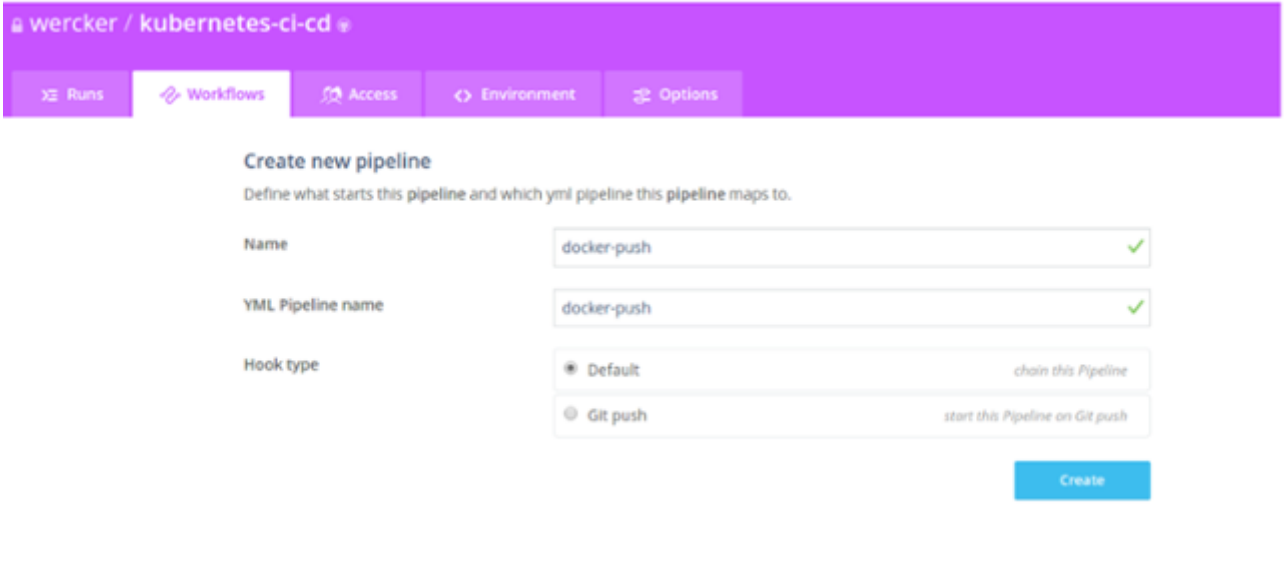
Configure Wercker.com GUI for this new pipeline:

- Go to the “Workflow” tab and click the “Add new pipeline” button.
- Name: docker-push.
- YML pipeline name: docker-push.



Configure Wercker.com GUI with our Quay.io credentials:

- Go to the “Environment” tab.
- Add a variable with key: QUAY\_USERNAME.
  - This is a username with permission to your quay.io repository.
- Add a variable with key: QUAY\_PASSWORD.
  - This is the password for the user.



2) Deploy the built container onto our Kubernetes cluster when there is a push into the `master` branch.

- a) Our definition for the “deploy” action: <https://github.com/wercker/kubernetes-ci-cd/blob/master/wercker.yml#L72>.

Configure Wercker.com GUI for this new pipeline:

- Go to the “Workflow” tab and click the “Add new pipeline” button.
- Name: deploy.
- YML pipeline name: deploy.

wercker / kubernetes-ci-cd

Runs Workflows Access Environment Options

**Create new pipeline**  
Define what starts this pipeline and which yml pipeline this pipeline maps to.

Name

YML Pipeline name

Hook type ☒ Default chain this Pipeline  
☐ Git push start this Pipeline on Git push

Create

Configure Wercker.com GUI with our Kubernetes information:

- Go to the “Environment” tab.
- Add a variable with key: KUBE\_ENDPOINT.
  - This is the URL to your Kubernetes API server. Remember that we gave Wercker.com’s IPs access to this above.
  - The `kops` tool created us a Kubernetes configuration file when we created the cluster. This is located in: `~/.kube/config`.
  - Run the command to output the entire config:
    - `cat ~/.kube/config.`
  - This will be the URL of the `server` key.
- Add a variable with key: KUBE\_USERNAME.
  - This is a username that Kubernetes uses for authentication.
  - You can find this in the `~/.kube/config` file.
  - There is a `username` key. This is the user. Unless you have changed it on build the user name is: admin.
- Add a variable with key: KUBE\_PASSWORD.
  - This is the password associated with the username.
  - This is the `password` key associated with the username.
- Add a variable with key: KUBE\_NAMESPACE.
  - Kubernetes allows you to segment off the cluster into what they call “namespaces”. For this tutorial we will use the default namespace called “default”.
  - The value for this is: default.

wercker / kubernetes-ci-cd

Runs

Workflows

Access

Environment

Options

Application environment variables

Settings and passwords defined here will be available to all pipelines

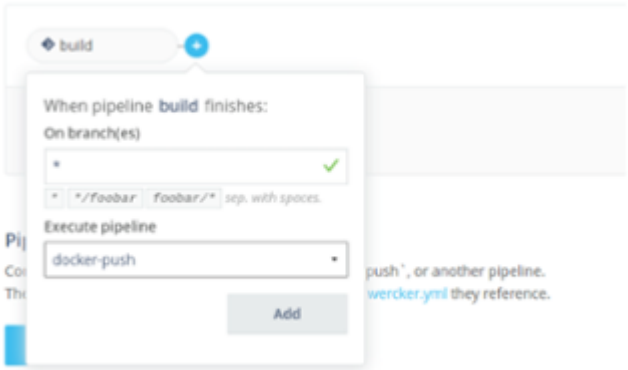
Key	Value	
QUAY_USERNAME	wercker+pusher	Delete
QUAY_PASSWORD	Protected	Delete
KUBE_ENDPOINT	https://api.dev1.k8s.squeich.io	Delete
KUBE_USERNAME	admin	Delete
KUBE_PASSWORD	Protected	Delete
KUBE_NAMESPACE	default	Delete

3) Creating the Pipeline. This strings along our “actions” above to the sequence we want which is:

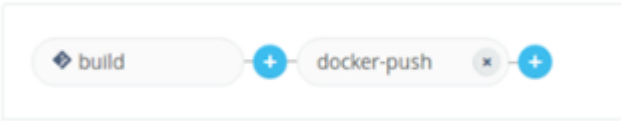
build -> push docker container -> deploy to our cluster if it is the master branch.

Configure Wercker.com GUI:

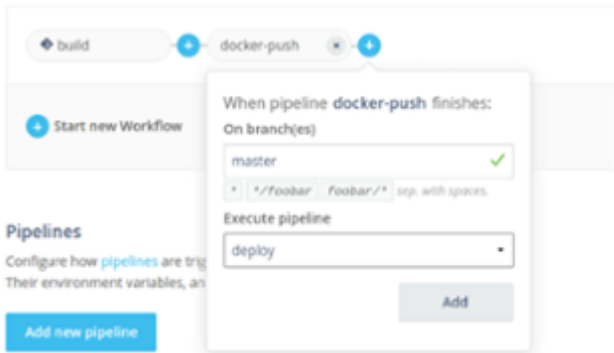
- Click on the “Workflow” tab.
- Click on the “+” on the right of the “build” icon.



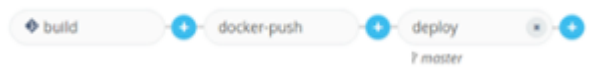
- In the “Select pipeline” dropbox select “docker-push”.
  - This tells our pipeline to run the docker-push action after a successful build



- Click on the “+” on the right of the “docker-push” action.
  - We will add the deploy next.
- This time we only want to invoke this action when it happens on the “master” branch.

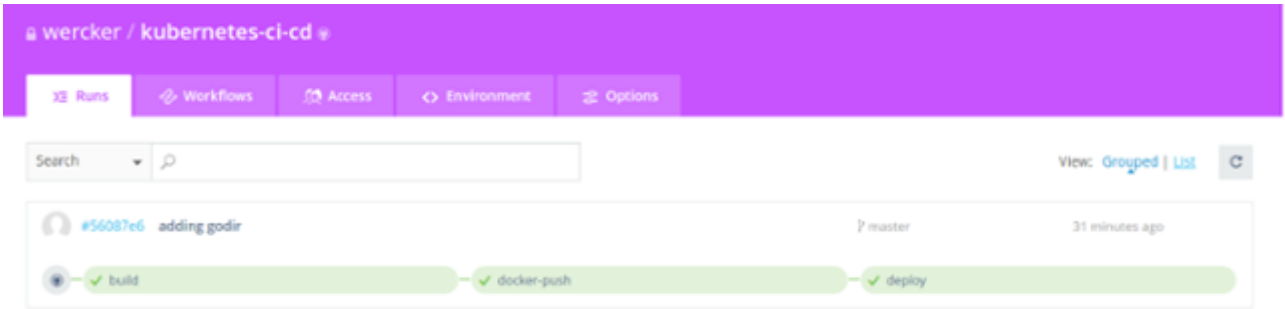


We have a “Workflow” that looks like this now. Notice the “master” under deploy which denotes this only happens for the “master” branch.



We have finished configuring Wercker.com to be able to deploy into our Kubernetes cluster.

To invoke it, we just have to make a small change in the Github repository and then push the code into master. You can do something as simple as changing/add something in the README.md and then pushing it in. You will get a pipeline run like this one:



You can check your Kubernetes cluster to see that this application has been deployed:

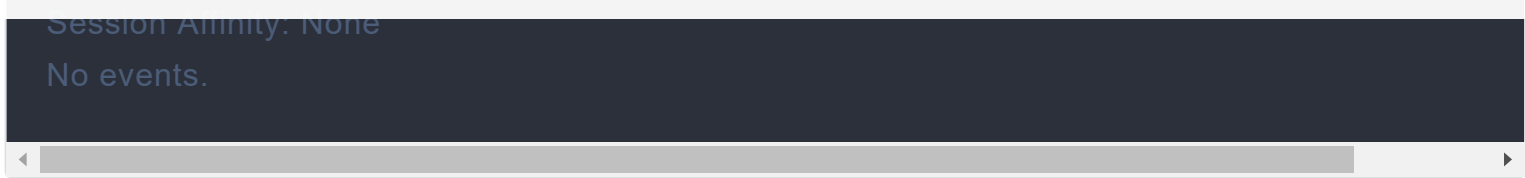
```
$ kubectl --namespace default get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP            NODE
ingress-controller-2171626210-5khg3  1/1     Running   0          13d   100.96.8.56   ip-172-31-17-103
ingress-default-backend-2249275363-1vnt6  1/1     Running   0          13d   100.96.8.57   ip-172-31-17-103
webapp-3028151003-5lb8f              1/1     Running   0          30m   100.96.11.43  ip-172-31-17-103
```

Check that the Kubernetes ingress has been deployed:

```
$ kubectl --namespace default get ing
NAME      HOSTS                                ADDRESS          PORTS   AGE
webapp    kubernetes-ci-cd-demo.wercker.com    54.172.67.240    80      49m
```

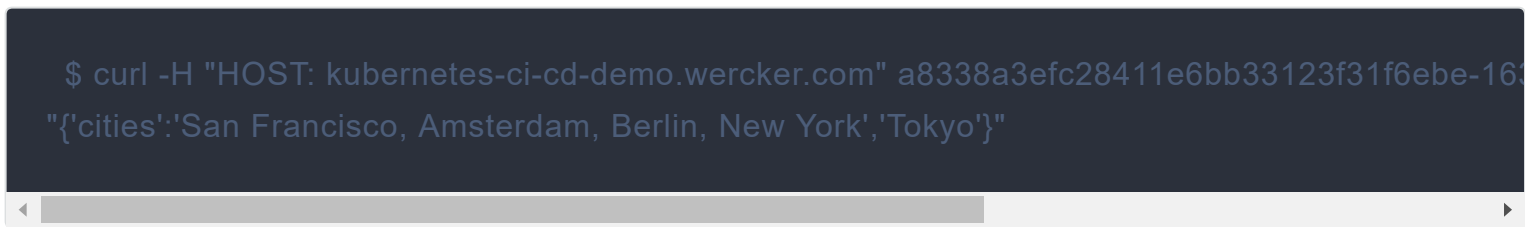
We can get the AWS ELB that Kubernetes created for us:

```
$ kubectl --namespace default describe svc ingress-lb
Name: ingress-lb
Namespace: default
Labels: <none>
Selector: app=ingress-controller
Type: LoadBalancer
IP: 100.69.179.81
LoadBalancer Ingress: a8338a3efc28411e6bb33123f31f6ebe-1633738224.us-east-1.elb.amazonaws.com
Port: http 80/TCP
NodePort: http 32660/TCP
Endpoints: 100.96.8.56:80
Port: https 443/TCP
NodePort: https 30841/TCP
Endpoints: 100.96.8.56:443
```



We can add a DNS CNAME for our Ingress to the AWS ELB address to reach it directly. Without the CNAME we can reach it via curl and adding in a host header.

The Kubernetes ingress which is running Nginx is doing virtual host routing:



## Conclusion

This tutorial walked you through setting up a Kubernetes cluster in a production configuration and then we walked through the setup of [wercker.com](https://wercker.com) to automatically build and deploy our application.

## Like Wercker?

We're hiring! Check out the [careers page](#) for open positions in Amsterdam, London and San Francisco.

As usual, if you want to stay in the loop follow us on twitter [@wercker](#) or hop on our [public slack channel](#). If it's your first time using Wercker, be sure to tweet out your `#greenbuilds`, and we'll send you some swag!

Topics: [Product](#), [Kubernetes](#), [Containers](#)