

dghubble Update Container Linux from 1576.5.0 to 1911.4.0

cd57013 13 days ago

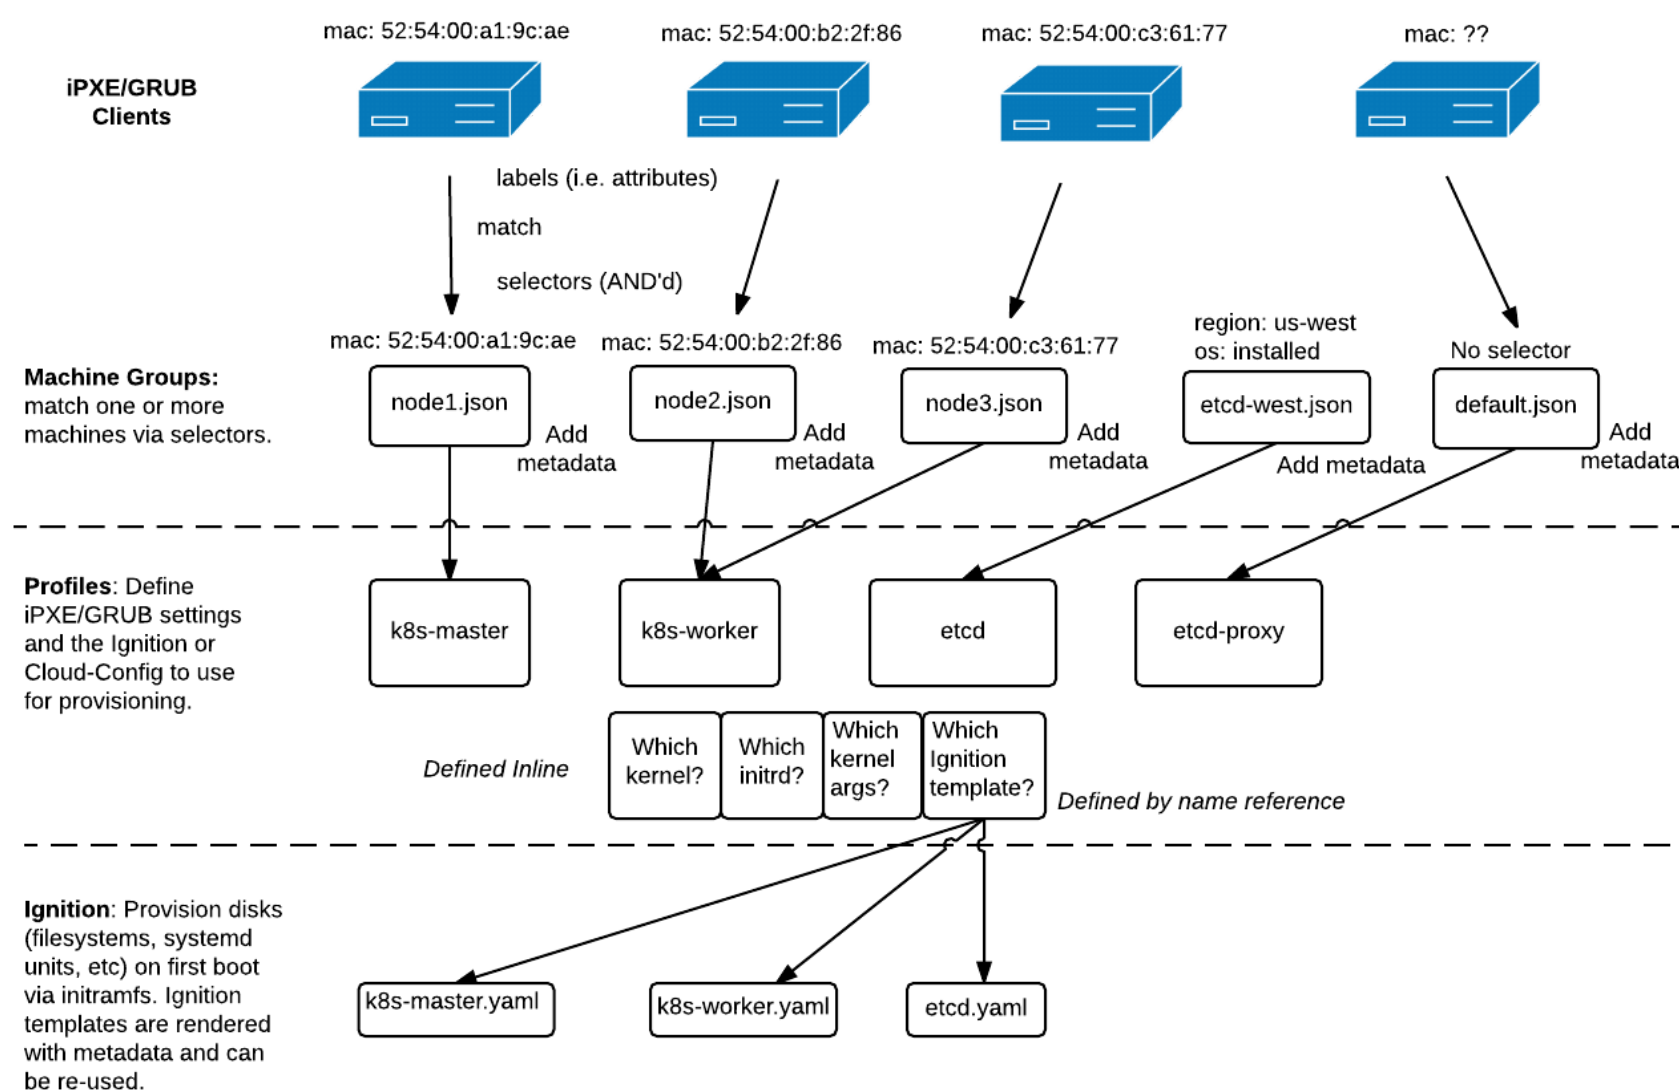
7 contributors

187 lines (138 sloc) 8.28 KB

matchbox

matchbox is an HTTP and gRPC service that renders signed [Ignition configs](#), [cloud-configs](#), network boot configs, and metadata to machines to create CoreOS Container Linux clusters. matchbox maintains **Group** definitions which match machines to *profiles* based on labels (e.g. MAC address, UUID, stage, region). A **Profile** is a named set of config templates (e.g. iPXE, GRUB, Ignition config, Cloud-Config, generic configs). The aim is to use Container Linux's early-boot capabilities to provision Container Linux machines.

Network boot endpoints provide PXE, iPXE, GRUB support. matchbox can be deployed as a binary, as an [appc](#) container with rkt, or as a Docker container.



Getting started

Get started running matchbox on your Linux machine, with rkt or Docker.

- [matchbox with rkt](#)
- [matchbox with Docker](#)

Flags

See [configuration](#) flags and variables.

API

- [HTTP API](#)
- [gRPC API](#)

Data

A `Store` stores machine Groups, Profiles, and associated Ignition configs, cloud-configs, and generic configs. By default, `matchbox` uses a `FileStore` to search a `-data-path` for these resources.

Prepare `/var/lib/matchbox` with `groups`, `profile`, `ignition`, `cloud`, and `generic` subdirectories. You may wish to keep these files under version control.

```

/var/lib/matchbox
├── cloud
│   ├── cloud.yaml.tpl
│   └── worker.sh.tpl
├── ignition
│   ├── raw.ign
│   ├── etcd.yaml.tpl
│   └── simple.yaml.tpl
├── generic
│   ├── config.yaml
│   ├── setup.cfg
│   └── datacenter-1.tpl
├── groups
│   ├── default.json
│   ├── node1.json
│   └── us-central1-a.json
└── profiles
    ├── etcd.json
    └── worker.json

```

The [examples](#) directory is a valid data directory with some pre-defined configs. Note that `examples/groups` contains many possible groups in nested directories for demo purposes (tutorials pick one to mount). Your machine groups should be kept directly inside the `groups` directory as shown above.

Profiles

Profiles reference an Ignition config, Cloud-Config, and/or generic config by name and define network boot settings.

```

{
  "id": "etcd",
  "name": "Container Linux with etcd2",
  "cloud_id": "",
  "ignition_id": "etcd.yaml",
  "generic_id": "some-service.cfg",
  "boot": {
    "kernel": "/assets/coreos/1911.4.0/coreos_production_pxe.vmlinuz",
    "initrd": ["/assets/coreos/1911.4.0/coreos_production_pxe_image.cpio.gz"],
    "args": [
      "coreos.config.url=http://matchbox.foo:8080/ignition?uuid=${uuid}&mac=${mac:hexhyp}",
      "coreos.first_boot=yes",
      "coreos.autologin"
    ]
  },
}

```

The `"boot"` settings will be used to render configs to network boot programs such as iPXE or GRUB. You may reference remote kernel and initrd assets or [local assets](#).

To use Ignition, set the `coreos.config.url` kernel option to reference the `matchbox` [Ignition endpoint](#), which will render the `ignition_id` file. Be sure to add the `coreos.first_boot` option as well.

To use cloud-config, set the `cloud-config-url` kernel option to reference the `matchbox` [Cloud-Config endpoint](#), which will render the `cloud_id` file.

Groups

Groups define selectors which match zero or more machines. Machine(s) matching a group will boot and provision according to the group's `Profile`.

Create a group definition with a `Profile` to be applied, selectors for matching machines, and any `metadata` needed to render templated configs. For example `/var/lib/matchbox/groups/node1.json` matches a single machine with MAC address `52:54:00:89:d8:10` .

```
# /var/lib/matchbox/groups/node1.json
{
  "name": "node1",
  "profile": "etcd",
  "selector": {
    "mac": "52:54:00:89:d8:10"
  },
  "metadata": {
    "fleet_metadata": "role=etcd,name=node1",
    "etcd_name": "node1",
    "etcd_initial_cluster": "node1=http://node1.example.com:2380,node2=http://node2.example.com:2380,node3=http://node3.example.com:2380"
  }
}
```

Meanwhile, `/var/lib/matchbox/groups/proxy.json` acts as the default machine group since it has no selectors.

```
{
  "name": "etcd-proxy",
  "profile": "etcd-proxy",
  "metadata": {
    "fleet_metadata": "role=etcd-proxy",
    "etcd_initial_cluster": "node1=http://node1.example.com:2380,node2=http://node2.example.com:2380,node3=http://node3.example.com:2380"
  }
}
```

For example, a request to `/ignition?mac=52:54:00:89:d8:10` would render the Ignition template in the "etcd" `Profile` , with the machine group's metadata. A request to `/ignition` would match the default group (which has no selectors) and render the Ignition in the "etcd-proxy" `Profile`. Avoid defining multiple default groups as resolution will not be deterministic.

Reserved selectors

Group selectors can use any key/value pairs you find useful. However, several labels have a defined purpose and will be normalized or parsed specially.

- `uuid` - machine UUID
- `mac` - network interface physical address (normalized MAC address)
- `hostname` - hostname reported by a network boot program
- `serial` - serial reported by a network boot program

Config templates

Profiles can reference various templated configs. Ignition JSON configs can be generated from [Container Linux Config](#) template files. Cloud-Config templates files can be used to render a script or Cloud-Config. Generic template files can be used to render arbitrary untyped configs (experimental). Each template may contain [Go template](#) elements which will be rendered with machine group metadata, selectors, and query params.

For details and examples:

- [Container Linux Config](#)
- [Cloud-Config](#)

Variables

Within Container Linux Config templates, Cloud-Config templates, or generic templates, you can use group metadata, selectors, or request-scoped query params. For example, a request `/generic?mac=52-54-00-89-d8-10&foo=some-param&bar=b` would match the `node1.json` machine group shown above. If the group's profile ("etcd") referenced a generic template, the following variables could be used.

```
# Untyped generic config file
# Selector
{{.mac}}                # 52:54:00:89:d8:10 (normalized)
# Metadata
```

```
{{.etcd_name}}           # node1
{{.fleet_metadata}}      # role=etcd,name=node1
# Query
{{.request.query.mac}}   # 52:54:00:89:d8:10 (normalized)
{{.request.query.foo}}   # some-param
{{.request.query.bar}}   # b
# Special Addition
{{.request.raw_query}}   # mac=52:54:00:89:d8:10&foo=some-param&bar=b
```

Note that `.request` is reserved for these purposes so group metadata with data nested under a top level "request" key will be overwritten.

Assets

`matchbox` can serve `-assets-path` static assets at `/assets`. This is helpful for reducing bandwidth usage when serving the kernel and initrd to network booted machines. The default `assets-path` is `/var/lib/matchbox/assets` or you can pass `-assets-path=""` to disable asset serving.

```
matchbox.foo/assets/
└─ coreos
   └─ VERSION
      ├── coreos_production_pxe.vmlinuz
      └─ coreos_production_pxe_image.cpio.gz
```

For example, a `Profile` might refer to a local asset `/assets/coreos/VERSION/coreos_production_pxe.vmlinuz` instead of `http://stable.release.core-os.net/amd64-usr/VERSION/coreos_production_pxe.vmlinuz`.

See the [get-coreos](#) script to quickly download, verify, and place Container Linux assets.

Network

`matchbox` does not implement or exec a DHCP/TFTP server. Read [network setup](#) or use the [coreos/dnsmasq](#) image if you need a quick DHCP, proxyDHCP, TFTP, or DNS setup.

Going further

- [gRPC API Usage](#)
- [Metadata](#)
- [OpenPGP Signing](#)