

March 18, 2017 · KUBERNETES

Self-hosted Kubernetes on bare-metal with Bootkube/Matchbox

10/22/2017: Updated post on this method

I use a cobbler VirtualBox VM on my laptop to PXE boot my three bare-metal servers in my home lab for OpenStack. This enables me to quickly test new OpenStack deployments with setting three "--netboot" cobbler values to true and then rebooting my servers. Cobbler takes care of PXE booting my servers with Ubuntu and with my specific partitioning scheme. I can then use Ansible to prepare my three nodes and then use Ansible to lay down OpenStack.

I was looking for a similar solution for testing Kubernetes when a friend at Rackspace pointed me to bootkube and matchbox. I've used a few methods for deploying a K8s cluster, a manually using Kelsey Hightowers **'Kubernetes the Hard Way'**, **minikube** for local use, and recently **kubeadm**, which was released with K8s 1.5. kubeadm makes setting up a test cluster extremely simple and if you're just wanting to quickly use Kubernetes in a multi-node setup, I would recommend kubeadm. I believe there will be some HA capabilities added to kubeadm in K8s 1.6 that I'm looking forward to trying out.

With bootkube and matchbox I was able to get a similar setup to cobbler where I can PXE boot my three bare-metal servers with **Container Linux** and then bootkube will bootstrap a self-hosted Kubernetes cluster. Matchbox uses groups/profiles and ignition configs. [This documentation](#) does a good job of describing Matchbox. It took me awhile to understand the general flow of this process.

Matchbox API -> Groups -> Profiles -> Ignition Configs

Instead of even attempting to explain what self-hosted means, I'll defer to **CoreOS** CTO Brandon Philips who wrote a great overview of self-hosted Kubernetes and how bootkube is attempting to solve this workflow.

<https://github.com/kubernetes/community/blob/master/contributors/design-proposals/self-hosted-kubernetes.md>

As you can see, deploying a self-hosted cluster enables some interesting cluster management abilities like managing the Kubernetes control plane just like any other application managed by Kubernetes.

Setup

Let's take a look at the matchbox setup. I spun up a VirtualBox VM on my laptop with Ubuntu 16.04 to use as my matchbox/bootkube/dnsmasq host. We'll need to create dnsmasq and matchbox containers. I did this setup with rkt instead of docker, but they have **documentation** for both implementations.

Install rkt

```
# git clone https://github.com/coreos/matchbox.git
```

This will download the Container Linux images locally so when it's served, it doesn't have to download those over the Internet.

```
# cd matchbox
# ./scripts/get-coreos stable 1235.9.0 ./examples/assets
```

Take note of the examples directory. It has some preconfigured examples for various things bootkube/matchbox can configure. I'm going to focus on one, **bootkube-install**. This group will install Container Linux to your servers and then bootstrap a self-hosted Kubernetes cluster.

```
root@bootkube:~/matchbox# ls examples/groups/bootkube-install
install.json  node1.json  node2.json  node3.json
```

```
root@bootkube:~/matchbox# cat examples/groups/bootkube-install/install.json
{
  "id": "coreos-install",
  "name": "CoreOS Install",
  "profile": "install-reboot",
  "metadata": {
    "coreos_channel": "stable",
    "coreos_version": "1235.9.0",
    "ignition_endpoint": "http://bootkube:8080/ignition",
    "baseurl": "http://bootkube:8080/assets/coreos"
  }
}
```

The bootkube hostname resolves to my VirtualBox VM running matchbox/bootkube/dnsmasq. This group will install Container Linux on the servers local disk from the 'install-reboot' profile. The node1-3 files are where we enter our MAC address for the bare metal servers and any SSH keys we want installed.

```
root@bootkube:~/matchbox# cat examples/groups/bootkube-install/node1.json
{
  "id": "node1",
  "name": "Controller Node",
  "profile": "bootkube-controller",
  "selector": {
    "mac": "44:39:C4:53:4C:39",
    "os": "installed"
  },
  "metadata": {
    "domain_name": "node1.example.com",
    "etcd_initial_cluster": "node1=http://node1:2380",
    "etcd_name": "node1",
    "k8s_dns_service_ip": "10.3.0.10",
    "ssh_authorized_keys": [
      "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDgFTTmM5K1IK9LkMzyk05/9dVe30AW037fecibe"
    ]
  }
}
```

node1 will act as the etcd/controller node. node2 is only a worker node.

```

root@bootkube:~/matchbox# cat examples/groups/bootkube-install/node2.json
{
  "id": "node2",
  "name": "Worker Node",
  "profile": "bootkube-worker",
  "selector": {
    "mac": "D4:AE:52:C8:A1:8D",
    "os": "installed"
  },
  "metadata": {
    "domain_name": "node2.example.com",
    "etcd_endpoints": "node1:2379",
    "k8s_dns_service_ip": "10.3.0.10",
    "ssh_authorized_keys": [
      "ssh-rsa AAAAB3NzaC1yc2EAAAADAQABAAQDgFTTmM5K1IK9LkMzyk05/9dVe30AW037fecibe
    ]
  }
}

```

Now let's take a look at the profiles. The first profile it uses is install-reboot.

```

root@bootkube:~/matchbox# ls examples/profiles
bootkube-controller.json  etcd3-gateway.json  grub.json  install-shutdown.json

root@bootkube:~/matchbox# cat examples/profiles/install-reboot.json
{
  "id": "install-reboot",
  "name": "Install CoreOS and Reboot",
  "boot": {
    "kernel": "/assets/coreos/1235.9.0/coreos_production_pxe.vmlinuz",
    "initrd": ["/assets/coreos/1235.9.0/coreos_production_pxe_image.cpio.gz"],
    "args": [
      "coreos.config.url=http://bootkube:8080/ignition?uuid=${uuid}&mac=${mac:hexhyp:}",
      "coreos.first_boot=yes",
      "console=tty0",
      "console=ttyS0",
      "coreos.autologin"
    ]
  },
  "ignition_id": "install-reboot.yaml"
}

```

The last line calls the ignition file that will be used.

```

root@bootkube:~/matchbox# cat examples/ignition/install-reboot.yaml
---
```

```

systemd:
  units:
    - name: installer.service
      enable: true
      contents: |
        [Unit]
        Requires=network-online.target
        After=network-online.target
        [Service]
        Type=simple
        ExecStart=/opt/installer
        [Install]
        WantedBy=multi-user.target

storage:
  files:
    - path: /opt/installer
      filesystem: root
      mode: 0500
      contents:
        inline: |
          #!/bin/bash -ex
          curl "{{.ignition_endpoint}}?{{.request.raw_query}}&os=installed" -o ignit:
          coreos-install -d /dev/sda -C {{.coreos_channel}} -V {{.coreos_version}} -:
          udevadm settle
          systemctl reboot

{{ if index . "ssh_authorized_keys" }}
passwd:
  users:
    - name: core
      ssh_authorized_keys:
        {{ range $element := .ssh_authorized_keys }}
        - {{$element}}
        {{end}}
  {{end}}

```

This is where it lays down the systemd unit files and other system tasks, like installing Container Linux to a specific device, "-d /dev/sda". You can customize or create new ignition files for a specific chassis, etc. Some docs can be found [here](#).

node1 then matches the bootkube-controller profile.

```

root@bootkube:~/matchbox# cat examples/profiles/bootkube-controller.json
{
  "id": "bootkube-controller",
  "name": "bootkube Ready Controller",
  "boot": {
    "kernel": "/assets/coreos/1235.9.0/coreos_production_pxe.vmlinuz",
    "initrd": ["/assets/coreos/1235.9.0/coreos_production_pxe_image.cpio.gz"],

```

```

    "args": [
      "root=/dev/sda1",
      "coreos.config.url=http://bootkube:8080/ignition?uuid=${uuid}&mac=${mac:hexhyp:}",
      "coreos.first_boot=yes",
      "console=tty0",
      "console=ttyS0",
      "coreos.autologin"
    ],
  },
  "ignition_id": "bootkube-controller.yaml"
}

```

The ignition file for the bootkube-controller.

```

root@bootkube:~/matchbox# cat examples/ignition/bootkube-controller.yaml
---
systemd:
  units:
    - name: etcd-member.service
      enable: true
      dropins:
        - name: 40-etcd-cluster.conf
          contents: |
            [Service]
            Environment="ETCD_IMAGE_TAG=v3.1.0"
            Environment="ETCD_NAME={{.etcd_name}}"
            Environment="ETCD_ADVERTISE_CLIENT_URLS=http://{{.domain_name}}:2379"
            Environment="ETCD_INITIAL_ADVERTISE_PEER_URLS=http://{{.domain_name}}:2380"
            Environment="ETCD_LISTEN_CLIENT_URLS=http://0.0.0.0:2379"
            Environment="ETCD_LISTEN_PEER_URLS=http://0.0.0.0:2380"
            Environment="ETCD_INITIAL_CLUSTER={{.etcd_initial_cluster}}"
            Environment="ETCD_STRICT_RECONFIG_CHECK=true"
        - name: docker.service
          enable: true
        - name: locksmithd.service
          dropins:
            - name: 40-etcd-lock.conf
              contents: |
                [Service]
                Environment="REBOOT_STRATEGY=etcd-lock"
        - name: kubelet.path
          enable: true
          contents: |
            [Unit]
            Description=Watch for kubeconfig
            [Path]
            PathExists=/etc/kubernetes/kubeconfig
            [Install]

```

```

    WantedBy=multi-user.target
- name: wait-for-dns.service
  enable: true
  contents: |
    [Unit]
    Description=Wait for DNS entries
    Wants=systemd-resolved.service
    Before=kubelet.service
    [Service]
    Type=oneshot
    RemainAfterExit=true
    ExecStart=/bin/sh -c 'while ! /usr/bin/grep '^[^#[:space:]]' /etc/resolv.conf
    [Install]
    RequiredBy=kubelet.service
- name: kubelet.service
  contents: |
    [Unit]
    Description=Kubelet via Hyperkube ACI
    [Service]
    Environment="RKT_OPTS=--uuid-file-save=/var/run/kubelet-pod.uuid \
      --volume=resolv,kind=host,source=/etc/resolv.conf \
      --mount volume=resolv,target=/etc/resolv.conf \
      --volume var-lib-cni,kind=host,source=/var/lib/cni \
      --mount volume=var-lib-cni,target=/var/lib/cni \
      --volume var-log,kind=host,source=/var/log \
      --mount volume=var-log,target=/var/log"
    EnvironmentFile=/etc/kubernetes/kubelet.env
    ExecStartPre=/bin/mkdir -p /etc/kubernetes/manifests
    ExecStartPre=/bin/mkdir -p /srv/kubernetes/manifests
    ExecStartPre=/bin/mkdir -p /etc/kubernetes/checkpoint-secrets
    ExecStartPre=/bin/mkdir -p /etc/kubernetes/cni/net.d
    ExecStartPre=/bin/mkdir -p /var/lib/cni
    ExecStartPre=-/usr/bin/rkt rm --uuid-file=/var/run/kubelet-pod.uuid
    ExecStart=/usr/lib/coreos/kubelet-wrapper \
      --kubeconfig=/etc/kubernetes/kubeconfig \
      --require-kubeconfig \
      --cni-conf-dir=/etc/kubernetes/cni/net.d \
      --network-plugin=cni \
      --lock-file=/var/run/lock/kubelet.lock \
      --exit-on-lock-contention \
      --pod-manifest-path=/etc/kubernetes/manifests \
      --allow-privileged \
      --hostname-override={{.domain_name}} \
      --node-labels=master=true \
      --cluster_dns={{.k8s_dns_service_ip}} \
      --cluster_domain=cluster.local
    ExecStop=-/usr/bin/rkt stop --uuid-file=/var/run/kubelet-pod.uuid
    Restart=always
    RestartSec=10
    [Install]

```

```

    WantedBy=multi-user.target
- name: bootkube.service
  contents: |
    [Unit]
    Description=Bootstrap a Kubernetes control plane with a temp api-server
    [Service]
    Type=simple
    WorkingDirectory=/opt/bootkube
    ExecStart=/opt/bootkube/bootkube-start
storage:
  {{ if index . "pxe" }}
  disks:
    - device: /dev/sda
      wipe_table: true
      partitions:
        - label: ROOT
  filesystems:
    - name: root
      mount:
        device: "/dev/sda1"
        format: "ext4"
        create:
          force: true
          options:
            - "-LRROOT"
  {{end}}
  files:
    - path: /etc/kubernetes/kubelet.env
      filesystem: root
      mode: 0644
      contents:
        inline: |
          KUBELET_ACI=quay.io/coreos/hyperkube
          KUBELET_VERSION=v1.5.2_coreos.2
    - path: /etc/hostname
      filesystem: root
      mode: 0644
      contents:
        inline:
          {{.domain_name}}
    - path: /etc/sysctl.d/max-user-watches.conf
      filesystem: root
      contents:
        inline: |
          fs.inotify.max_user_watches=16184
    - path: /opt/bootkube/bootkube-start
      filesystem: root
      mode: 0544
      user:
        id: 500

```

```

group:
  id: 500
contents:
  inline: |
    #!/bin/bash
    # Wrapper for bootkube start
    set -e
    BOOTKUBE_ACI="${BOOTKUBE_ACI:-quay.io/coreos/bootkube}"
    BOOTKUBE_VERSION="${BOOTKUBE_VERSION:-v0.3.7}"
    BOOTKUBE_ASSETS="${BOOTKUBE_ASSETS:-/opt/bootkube/assets}"
    exec /usr/bin/rkt run \
      --trust-keys-from-https \
      --volume assets,kind=host,source=$BOOTKUBE_ASSETS \
      --mount volume=assets,target=/assets \
      $RKT_OPTS \
      ${BOOTKUBE_ACI}:${BOOTKUBE_VERSION} --net=host --exec=/bootkube -- start

{{ if index . "ssh_authorized_keys" }}
passwd:
  users:
    - name: core
      ssh_authorized_keys:
        {{ range $element := .ssh_authorized_keys }}
        - {{ $element }}
        {{ end }}
  {{ end }}

```

It lays down some more systemd unit files and other system things. It sets up the service that monitors for "/etc/kubernetes/kubeconfig", which will then kick off kubelet via hyperkube. It also creates /opt/bootkube/bootkube-start, which is a wrapper for the temporary bootkube rkt container.

Install bootkube on the VirtualBox VM so we can generate all the asset files our hosts require.

```

root@bootkube:~/matchbox# wget https://github.com/kubernetes-incubator/bootkube/releases
root@bootkube:~/matchbox# tar xzf bootkube.tar.gz
root@bootkube:~/matchbox# ./bin/linux/bootkube version
Version: v0.3.9

```

```

root@bootkube:~/matchbox# ./bin/linux/bootkube render --asset-dir=assets --api-server=

```

That should generate files in assets/

```

root@bootkube:~/matchbox# ls assets/
auth  manifests  tls

```

The manifests files are what are used to build the self-hosted Kubernetes control plane.


```
root@bootkube:~/matchbox# ls assets/manifests/
kube-apiserver-secret.yaml  kube-controller-manager-disruption.yaml  kube-controller-
```

Now we're ready to start the matchbox/dnsmasq containers. These are values that worked for my setup, you might have to change these up a bit. If you wanted to create a new group/profile/ignition directory structure, you can do that and mount your own files in these commands.

```
root@bootkube:~/matchbox# rkt run coreos.com/dnsmasq:v0.3.0 --net=host -- -d -q --dhc
```

```
root@bootkube:~/matchbox# rkt run --net=host --mount volume=data,target=/var/lib/matchbox
```

When starting the matchbox container, we specify the "source=\$PWD/examples/groups/bootkube-install" location for the files we edited.

At this point you should be able to reboot your servers, they will install CoreOS, then reboot and then run the bootkube-controller and bootkube-worker profiles, which run the ignition files for those profiles.

Kubernetes cluster setup

The first action your controller node will attempt is to start up an etcd cluster.

```
node1 ~ # systemctl status etcd-member
● etcd-member.service - etcd (System Application Container)
   Loaded: loaded (/usr/lib/systemd/system/etcd-member.service; enabled; vendor prese
  Drop-In: /etc/systemd/system/etcd-member.service.d
           └─40-etcd-cluster.conf
   Active: active (running) since Thu 2017-03-16 04:35:13 UTC; 5min ago
     Docs: https://github.com/coreos/etcd
  Process: 7051 ExecStartPre=/usr/bin/rkt rm --uuid-file=/var/lib/coreos/etcd-member-
  Process: 7024 ExecStartPre=/usr/bin/mkdir --parents /var/lib/coreos (code=exited, :
 Main PID: 7090 (etcd)
    Tasks: 10
   Memory: 27.5M
      CPU: 1.298s
   CGroup: /system.slice/etcd-member.service
           └─7090 /usr/local/bin/etcd
```

```
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.554742 I | raft: a9aee6
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.554752 I | raft: a9aee6
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.554757 I | raft: raft.i
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.554919 I | etcdserver:
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.554945 I | etcdserver:
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.554953 I | embed: ready
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.555170 N | embed: serv:
Mar 16 04:35:13 node1 systemd[1]: Started etcd (System Application Container).
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.556941 N | etcdserver/r
Mar 16 04:35:13 node1 etcd-wrapper[7090]: 2017-03-16 04:35:13.556969 I | etcdserver/;
```

So now we have a working etcd cluster for Kubernetes.

```
node1 ~ # etcdctl cluster-health
member a9aee06e6a14d468 is healthy: got healthy result from http://node1:2379
cluster is healthy
```

Secure copy the kubeconfig to /etc/kubernetes/kubeconfig on every node which will activate kubelet.service. The /etc/systemd/system/kubelet.service will start kubelet via hyperkube ACI.

```
root@bootkube:~/matchbox# scp assets/auth/kubeconfig core@node1:/home/core/kubeconfig
root@bootkube:~/matchbox# ssh core@node1 'sudo mv kubeconfig /etc/kubernetes/kubecon-
```

Now we move the asset files over and start up bootkube, which will setup our temporary API and bootstrap a self-hosted Kubernetes cluster.

```
root@bootkube:~/matchbox# scp -r assets core@node1:/home/core
root@bootkube:~/matchbox# ssh core@node1 'sudo mv assets /opt/bootkube/assets && sudo
```

You can view this by tailing the bootkube services log. Once it's done, you should see the following.

```
node1 ~ # journalctl -f -u bootkube

...

Mar 17 21:59:36 node1 bootkube-start[8093]: [149795.171672] bootkube[5]: All self-ho-
```

Here's rkt and docker output. As you can see, the temporary bootkube container has already stopped.

```
node1 ~ # rkt list
UUID          APP      IMAGE NAME                                STATE  CREATED      STARTED      NETW
615573b4      etcd     quay.io/coreos/etcd:v3.1.0               running 1 day ago   1 day
75cb7ba5      hyperkube quay.io/coreos/hyperkube:v1.5.2_coreos.2 running 44 minutes
a6848aa8      bootkube quay.io/coreos/bootkube:v0.3.7            exited 8 minutes ago

node1 ~ # docker ps
CONTAINER ID      IMAGE
8aec69372861      quay.io/coreos/hyperkube:v1.5.2_coreos.2
285099a245ec      quay.io/coreos/hyperkube:v1.5.2_coreos.2
1027ddd38138      quay.io/coreos/hyperkube:v1.5.2_coreos.2
76b515b653ed      quay.io/coreos/hyperkube:v1.5.2_coreos.2
54cb191700fc      quay.io/coreos/hyperkube:v1.5.2_coreos.2
8f1ab95709c5      gcr.io/google_containers/pause-amd64:3.0
```

12/16/2018Self-hosted Kubernetes on bare-metal with Bootkube/Matchbox

```
37c55a435736      gcr.io/google_containers/pause-amd64:3.0
6351298ade64      gcr.io/google_containers/pause-amd64:3.0
fa1098928dda      gcr.io/google_containers/pause-amd64:3.0
506ef2593ecb      busybox
e589d8122961      quay.io/coreos/pod-checkpointer:5b585a2d731173713fa6871c436f6c53-
c30794abbc0a      quay.io/coreos/flannel:v0.7.0-amd64
5ca78d87f35c      gcr.io/google_containers/pause-amd64:3.0
b129fa7d747d      quay.io/coreos/pod-checkpointer:5b585a2d731173713fa6871c436f6c53-
4b341ce04e34      quay.io/coreos/hyperkube:v1.5.2_coreos.2
04a038ef35fe      gcr.io/google_containers/pause-amd64:3.0
7473cfa44f71      gcr.io/google_containers/pause-amd64:3.0
febfefa18eee      gcr.io/google_containers/pause-amd64:3.0
0d7d3db10141      gcr.io/google_containers/pause-amd64:3.0
```

For your worker nodes, you only need to copy over the kubeconfig file and it will activate the services needed and will automatically join the Kubernetes cluster, assuming your kubeconfig is correct.

Install the kubectl binary on your machine and copy over the kubeconfig file from assets/auth/kubeconfig. Then on your local machine, set the the KUBECONFIG variable to your kubeconfig file and check out your Kubernetes cluster!

```
export KUBECONFIG=kubeconfig
```

```
[17:09]shane@work~$ kubectl cluster-info
Kubernetes master is running at https://node1:443
KubeDNS is running at https://node1:443/api/v1/proxy/namespaces/kube-system/services,

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

Woot!


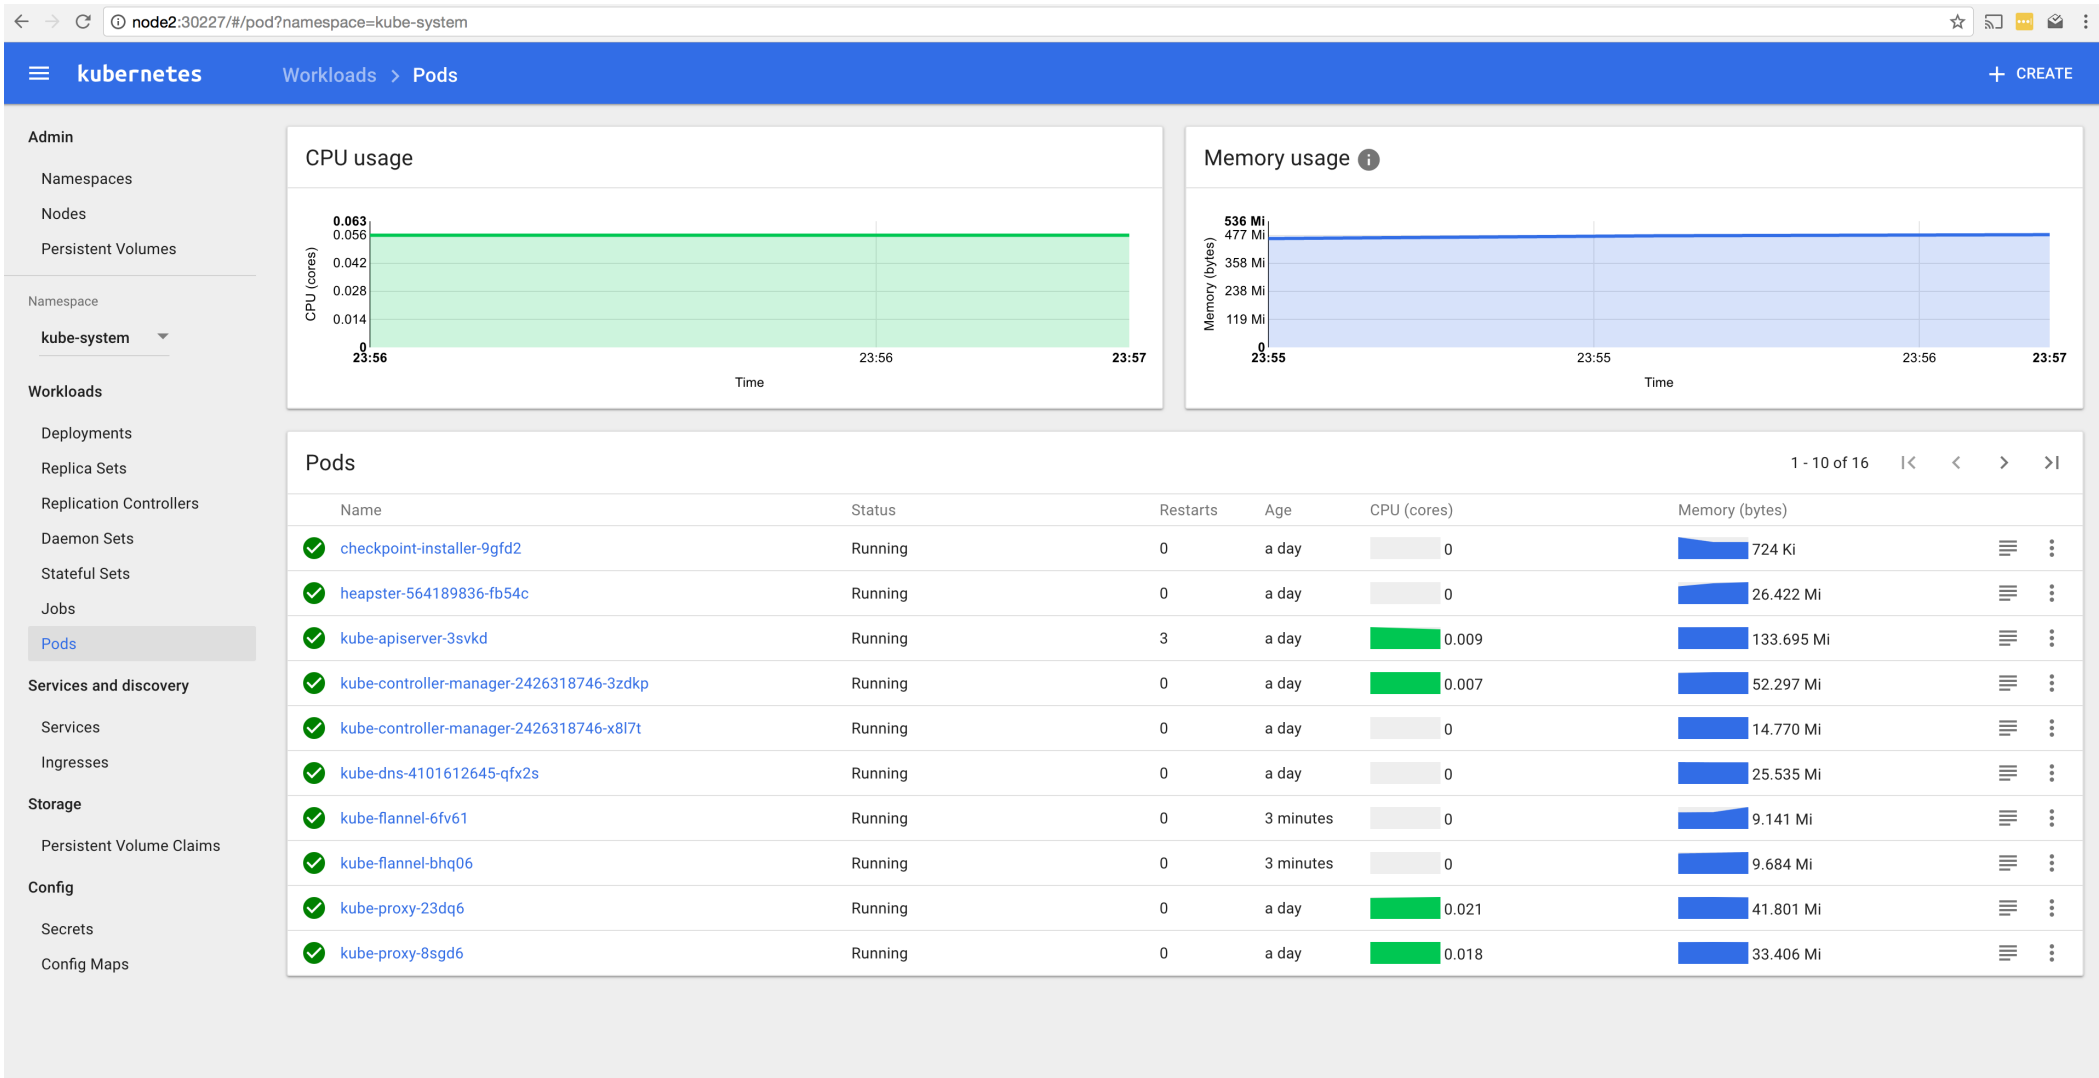
```
[17:12]shane@work~$ kubectl get nodes
NAME      STATUS    AGE
node1     Ready     14m
node2     Ready     14m

[17:12]shane@work~$ kubectl get pods --all-namespaces
NAMESPACE      NAME                                                    READY    STATUS    RESTARTS
kube-system    checkpoint-installer-fpdvj                            1/1      Running   0
kube-system    kube-apiserver-gtjmv                                    1/1      Running   2
kube-system    kube-controller-manager-2426318746-bbn1k              1/1      Running   0
kube-system    kube-controller-manager-2426318746-v84jr              1/1      Running   0
kube-system    kube-dns-4101612645-j745c                             4/4      Running   0
kube-system    kube-flannel-8vhwn                                     2/2      Running   0
kube-system    kube-flannel-pzv0l                                     2/2      Running   1
kube-system    kube-proxy-w3qqh                                       1/1      Running   0
kube-system    kube-proxy-zjxq1                                       1/1      Running   0
kube-system    kube-scheduler-2947727816-3z5mq                       1/1      Running   0
```

12/16/2018Self-hosted Kubernetes on bare-metal with Bootkube/Matchbox

kube-system	kube-scheduler-2947727816-j2lzl	1/1	Running	0
kube-system	pod-checkpointer-node1	1/1	Running	0

Here are some screenshots from the Kubernetes dashboard from a two node setup.



Not knowing about bootkube and matchbox two weeks ago I'm really impressed with the work that's gone into the projects. It makes the full process of describing your environment in YAML/JSON files (something I like from openstack-ansible) to PXE booting and bootstrapping bare metal servers into a Kubernetes cluster impressively simple. All of these manual steps can easily be converted to an Ansible playbook or any other configuration management software. The projects are moving quickly and some of this workflow will probably change. I'm looking forward to where these projects

