

[Back to Blog](#)

# MySQL on Docker: Deploy a Homogeneous Galera Cluster with etcd


Ashraf Sharif

December 07, 2016

Posted in:

[DB Ops Devops MySQL on Docker](#)

In the previous blog post, we have looked into the [multi-host networking capabilities with Docker with native network and Calico](#). In this blog post, our journey to make Galera Cluster run smoothly on Docker containers continues. Deploying Galera Cluster on Docker is tricky when using orchestration tools. Due to the nature of the scheduler in container orchestration tools and the assumption of homogenous images, the scheduler will just fire the respective containers according to the run command and leave the bootstrapping process to the container's entrypoint logic when starting up. And you do not want to do that for Galera - starting all nodes at once means each node will form a "1-node cluster" and you'll end up with a disjointed system.



## “Homogeneousing” Galera Cluster

Related posts

- [MySQL on Docker: Multi-Host Networking for MySQL Containers \(Part 2 - Calico\)](#)
- [MySQL on Docker: Introduction to Docker Swarm Mode and Multi-Host Networking](#)
- [MySQL on Docker: Building the Container Image](#)
- [MySQL Docker Containers: Understanding the basics](#)

That might be a new word, but it holds true for stateful services like MySQL Replication and Galera Cluster. As one might know, the bootstrapping process for Galera Cluster usually requires manual intervention, where you usually have to decide which node is the most advanced node to start bootstrapping from. There is nothing wrong with this step, you need to be aware of the state of each database node before deciding on the sequence of how to start them up. Galera Cluster is a distributed system, and its redundancy model works like that.

However, container orchestration tools like Docker Engine Swarm Mode and Kubernetes are not aware of the redundancy model of Galera. The orchestration tool presumes containers are independent from each other. If they are dependent, then you have to have an external service that monitors the state. The best way to achieve this is to use a key/value store as a reference point for other containers when starting up.

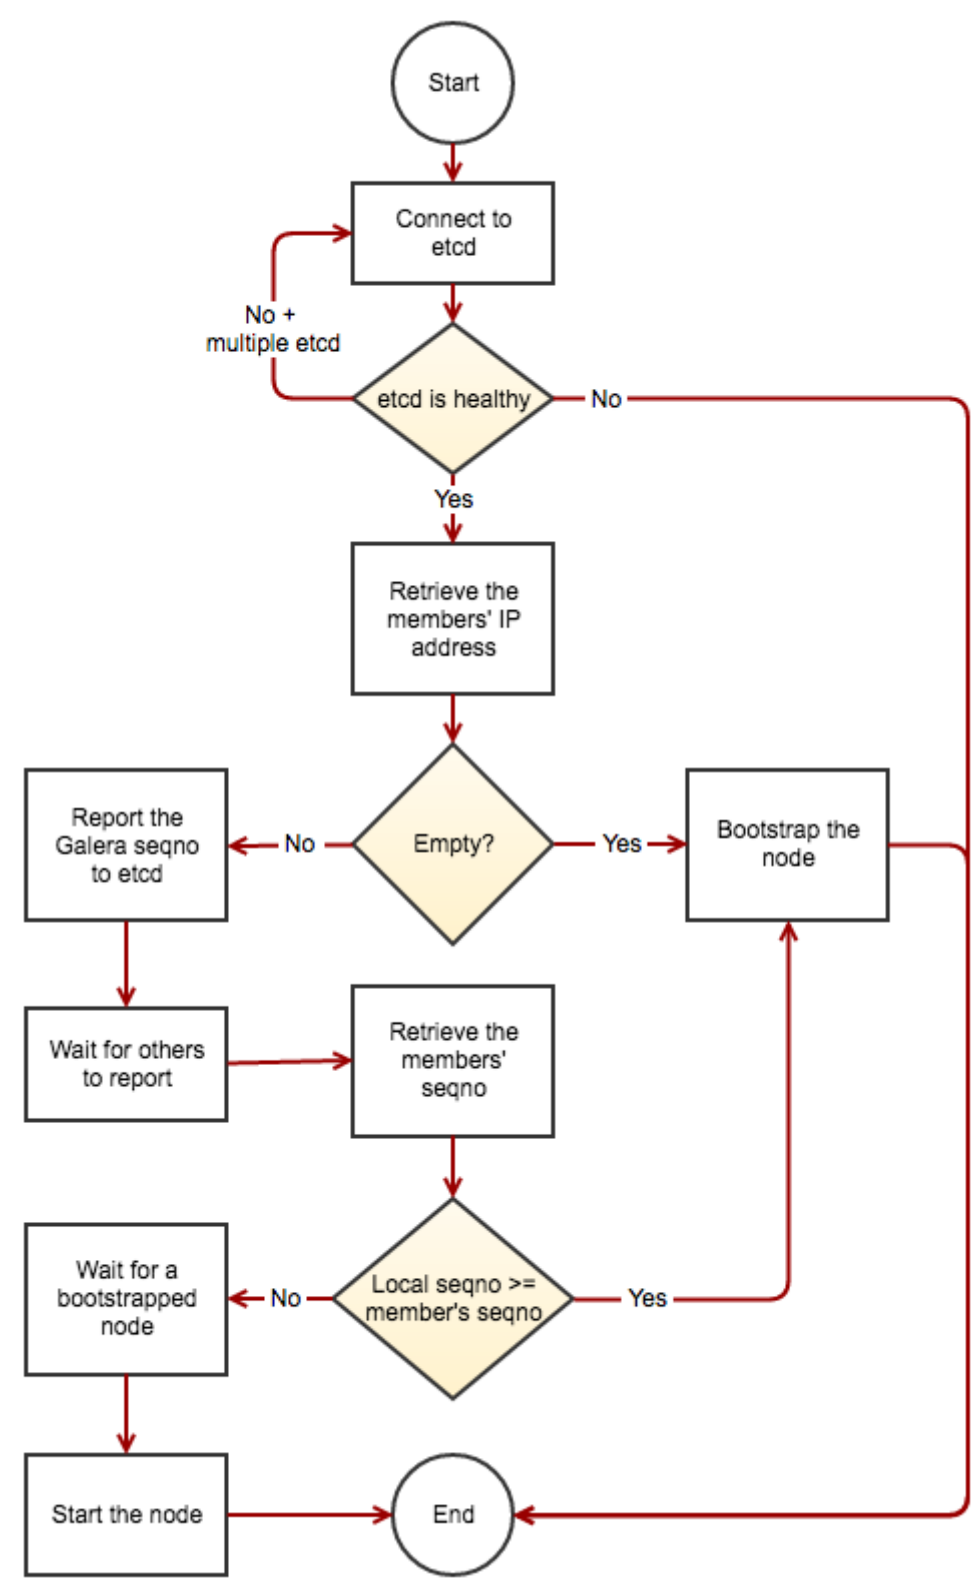
This is where service discovery like etcd comes into the picture. The basic idea is, each node should report its state periodically to the service. This simplifies the decision process when starting up. For Galera Cluster, the node that has `wsrep_local_state_comment` equal to `Synced` shall be used as a reference node when constructing the Galera communication address (`gcomm`) during joining. Otherwise, the most updated node has to get bootstrapped first.

Etcd has a very nice feature called TTL, where you can expire a key after a certain amount of time. This is useful to determine the state of a node, where the key/value entry only exists if an alive node reports to it. As a result, the node won't have to connect to each other to determine state (which is very troublesome in a dynamic environment) when forming a cluster. For example, consider the following keys:

```
1  {
2    "createdIndex": 10074,
3    "expiration": "2016-11-29T10:55:35.218496083Z",
4    "key": "/galera/my_wsrep_cluster/10.255.0.7/wsrep_last_committed",
5    "modifiedIndex": 10074,
6    "ttl": 10,
7    "value": "2881"
8  },
9  {
10   "createdIndex": 10072,
11   "expiration": "2016-11-29T10:55:34.650574629Z",
12   "key": "/galera/my_wsrep_cluster/10.255.0.7/wsrep_local_state_comment",
13   "modifiedIndex": 10072,
14   "ttl": 10,
15   "value": "Synced"
16 }
```

After 10 seconds (ttl value), those keys will be removed from the entry. Basically, all nodes should report to etcd periodically with an expiring key. Container should report every N seconds when it's alive (`wsrep_cluster_state_comment=Synced` and `wsrep_last_committed=#value`) via a background process. If a container is down, it will no longer send the update to etcd, thus the keys are removed after expiration. This simply indicates that the node was registered but is no longer synced with the cluster. It will be skipped when constructing the Galera communication address at a later point.

The overall flow of joining procedure is illustrated in the following flow chart:



We have built a Docker image that follows the above. It is specifically built for running Galera Cluster using Docker’s orchestration tool. It is available at [Docker Hub](#) and [our Github repository](#). It requires an etcd cluster as the discovery service (supports multiple etcd hosts) and based on Percona XtraDB Cluster 5.6. The image includes Percona Xtrabackup, jq (JSON processor) and also a shell script tailored for Galera health check called report\_status.sh.

You are welcome to fork or contribute to the project. Any bugs can be reported via [Github](#) or via our [support page](#).

## Deploying etcd Cluster

etcd is a distributed key value store that provides a simple and efficient way to store data across a cluster of machines. It’s open-source and available on GitHub. It provides shared configuration and service discovery. A simple use-case is to store database connection details or feature flags in etcd as key value pairs. It gracefully handles leader elections during network partitions and will tolerate machine failures, including the leader.

Since etcd is the brain of the setup, we are going to deploy it as a cluster daemon, on three nodes, instead of using containers. In this example, we are going to install etcd on each of the Docker hosts and form a three-node etcd cluster for better availability.

We used CentOS 7 as the operating system, with Docker v1.12.3, build 6b644ec. The deployment steps in this blog post are basically similar to the one used in our previous blog post.

1. Install etcd packages:

```
1 | $ yum install etcd
```

2. Modify the configuration file accordingly depending on the Docker hosts:

```
1 | $ vim /etc/etcd/etcd.conf
```

For docker1 with IP address 192.168.55.111:

```
1 | ETCD_NAME=etcd1
2 | ETCD_DATA_DIR="/var/lib/etcd/default.etcd"
3 | ETCD_LISTEN_PEER_URLS="http://0.0.0.0:2380"
4 | ETCD_LISTEN_CLIENT_URLS="http://0.0.0.0:2379"
5 | ETCD_INITIAL_ADVERTISE_PEER_URLS="http://192.168.55.111:2380"
6 | ETCD_INITIAL_CLUSTER="etcd1=http://192.168.55.111:2380,etcd2=http://192.168.55.112:2380,etcd3=http://192.168.55.113:2380"
7 | ETCD_INITIAL_CLUSTER_STATE="new"
8 | ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster-1"
9 | ETCD_ADVERTISE_CLIENT_URLS="http://0.0.0.0:2379"
```

For docker2 with IP address 192.168.55.112:

```
1 ETCD_NAME=etcd2
2 ETCD_DATA_DIR="/var/lib/etcd/default.etcd"
3 ETCD_LISTEN_PEER_URLS="http://0.0.0.0:2380"
4 ETCD_LISTEN_CLIENT_URLS="http://0.0.0.0:2379"
5 ETCD_INITIAL_ADVERTISE_PEER_URLS="http://192.168.55.112:2380"
6 ETCD_INITIAL_CLUSTER="etcd1=http://192.168.55.111:2380,etcd2=http://192.168.55.112:2380,etcd3=http://192.168.55.113:2380"
7 ETCD_INITIAL_CLUSTER_STATE="new"
8 ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster-1"
9 ETCD_ADVERTISE_CLIENT_URLS="http://0.0.0.0:2379"
```

For docker3 with IP address 192.168.55.113:

```
1 ETCD_NAME=etcd3
2 ETCD_DATA_DIR="/var/lib/etcd/default.etcd"
3 ETCD_LISTEN_PEER_URLS="http://0.0.0.0:2380"
4 ETCD_LISTEN_CLIENT_URLS="http://0.0.0.0:2379"
5 ETCD_INITIAL_ADVERTISE_PEER_URLS="http://192.168.55.113:2380"
6 ETCD_INITIAL_CLUSTER="etcd1=http://192.168.55.111:2380,etcd2=http://192.168.55.112:2380,etcd3=http://192.168.55.113:2380"
7 ETCD_INITIAL_CLUSTER_STATE="new"
8 ETCD_INITIAL_CLUSTER_TOKEN="etcd-cluster-1"
9 ETCD_ADVERTISE_CLIENT_URLS="http://0.0.0.0:2379"
```

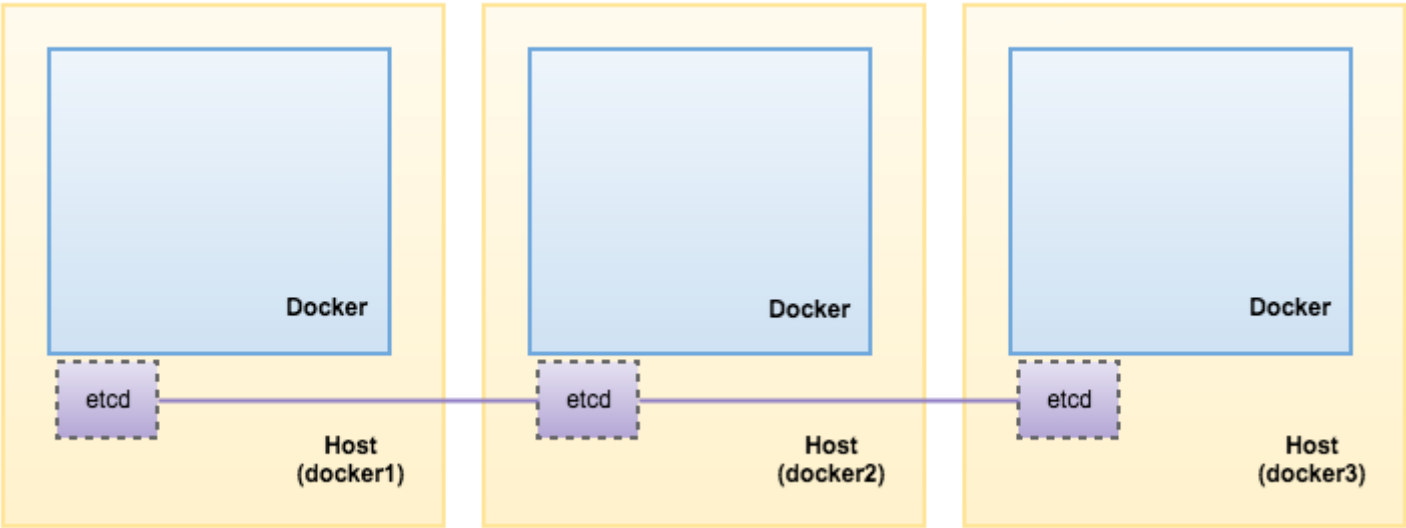
3. Start the service on docker1, followed by docker2 and docker3:

```
1 $ systemctl enable etcd
2 $ systemctl start etcd
```

4. Verify our cluster status using etcdctl:

```
1 [docker3 ]$ etcdctl cluster-health
2 member 2f8ec0a21c11c189 is healthy: got healthy result from http://0.0.0.0:2379
3 member 589a7883a7ee56ec is healthy: got healthy result from http://0.0.0.0:2379
4 member fcacfa3f23575abe is healthy: got healthy result from http://0.0.0.0:2379
5 cluster is healthy
```

That’s it. Our etcd is now running as a cluster on three nodes. The below illustrates our architecture:



## Deploying Galera Cluster

Minimum of 3 containers is recommended for high availability setup. Thus, we are going to create 3 replicas to start with, it can be scaled up and down afterwards. Running standalone is also possible with standard "docker run" command as shown further down.

Before we start, it’s a good idea to remove any sort of keys related to our cluster name in etcd:

```
1 $ etcdctl rm /galera/my_wsrep_cluster --recursive
```

### Ephemeral Storage

This is a recommended way if you plan on scaling the cluster out on more nodes (or scale back by removing nodes). To create a three-node Galera Cluster with ephemeral storage (MySQL datadir will be lost if the container is removed), you can use the following command:

```
1 $ docker service create \
2 --name mysql-galera \
3 --replicas 3 \
4 -p 3306:3306 \
5 --network galera-net \
6 --env MYSQL_ROOT_PASSWORD=mypassword \
7 --env DISCOVERY_SERVICE=192.168.55.111:2379,192.168.55.112:2379,192.168.55.113:2379 \
8 --env XTRABACKUP_PASSWORD=mypassword \
9 --env CLUSTER_NAME=my_wsrep_cluster \
10 severalnines/pxc56
```

### Persistent Storage

To create a three-node Galera Cluster with persistent storage (MySQL datadir persists if the container is removed), add the mount option with *type=volume*:

```
1 $ docker service create \
2 --name mysql-galera \
3 --replicas 3 \
4 -p 3306:3306 \
5 --network galera-net \
```

